

Convolutional Neural Network

Construction of Standard CNNs

Quang-Vinh Dinh
Ph.D. in Computer Science

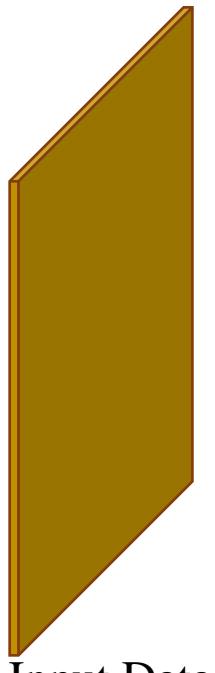
Outline

- Fashion-MNIST/Cifar10 Using only Conv2d
- Pooling
- Padding
- 1x1 Convolution
- LeNet and VGG Models

Convolutional Neural Network

❖ Convolution layer in PyTorch

```
nn.Conv2d(in_channels, out_channels, kernel_size)  
nn.ReLU()
```



Convolve with
1 kernel (1,5,5)
+activation



Convolve with
1 kernel (3,5,5)
+activation



Convolutional Neural Network

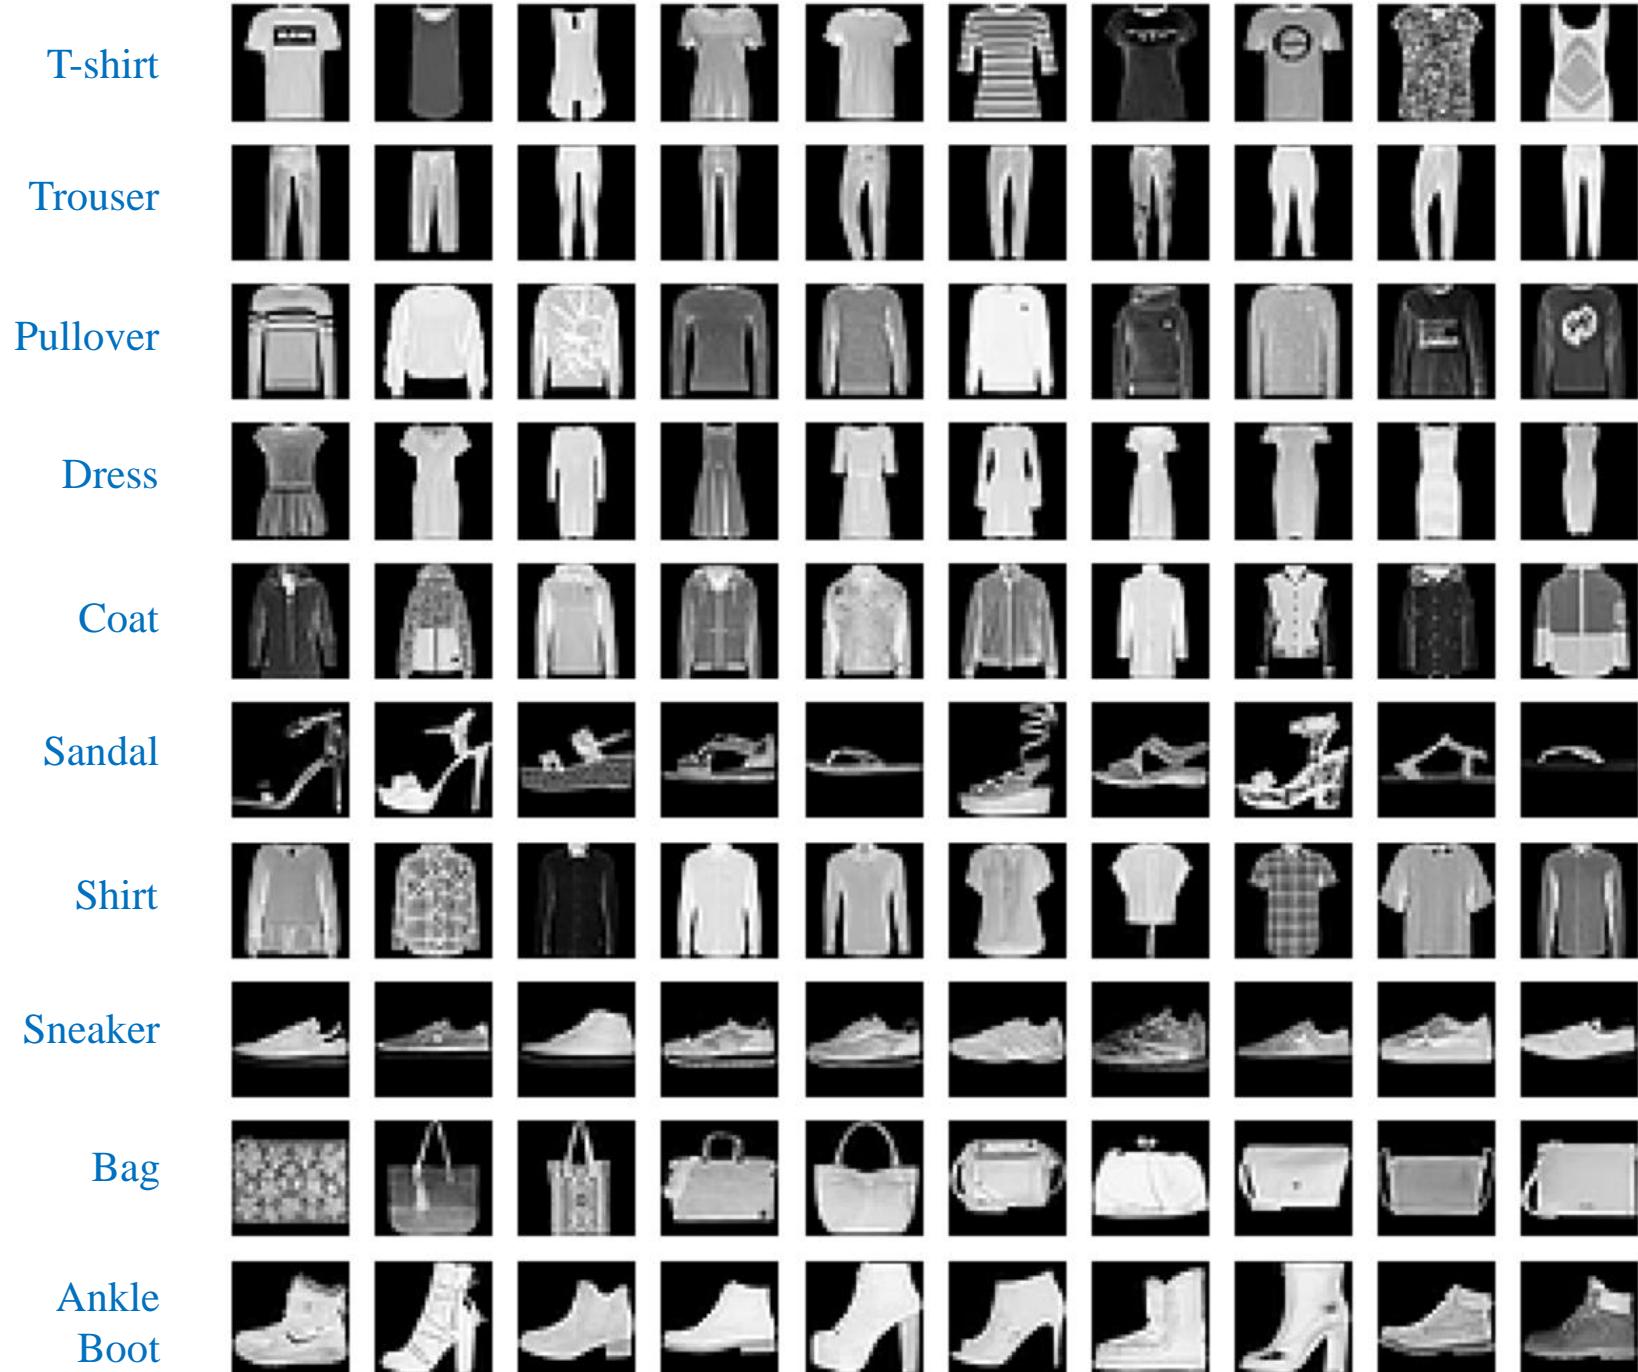
Fashion-MNIST dataset

Grayscale images

Resolution=28x28

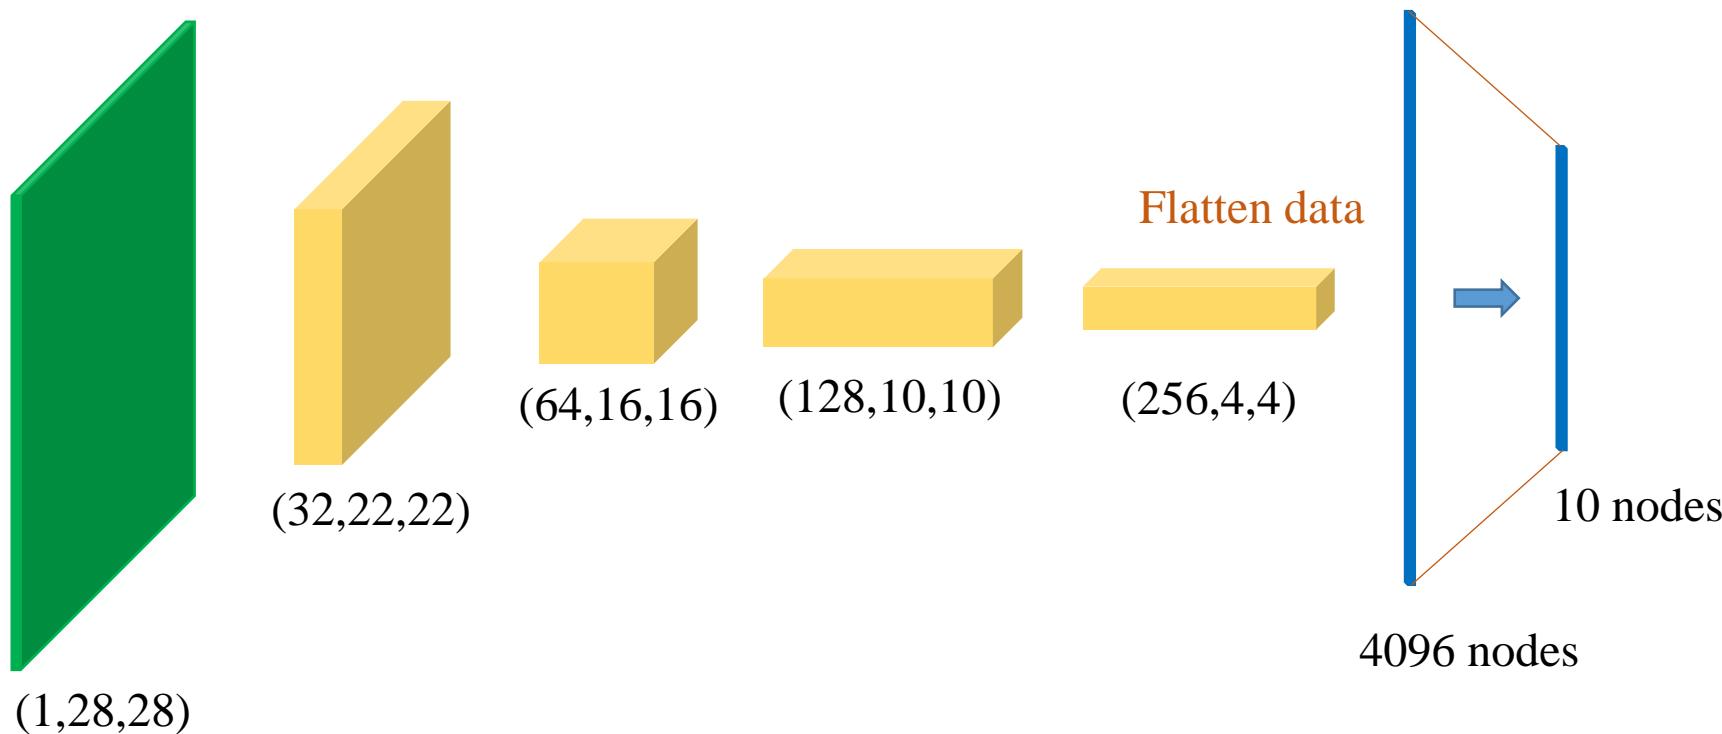
Training set: 60000 samples

Testing set: 10000 samples



Convolutional Neural Network

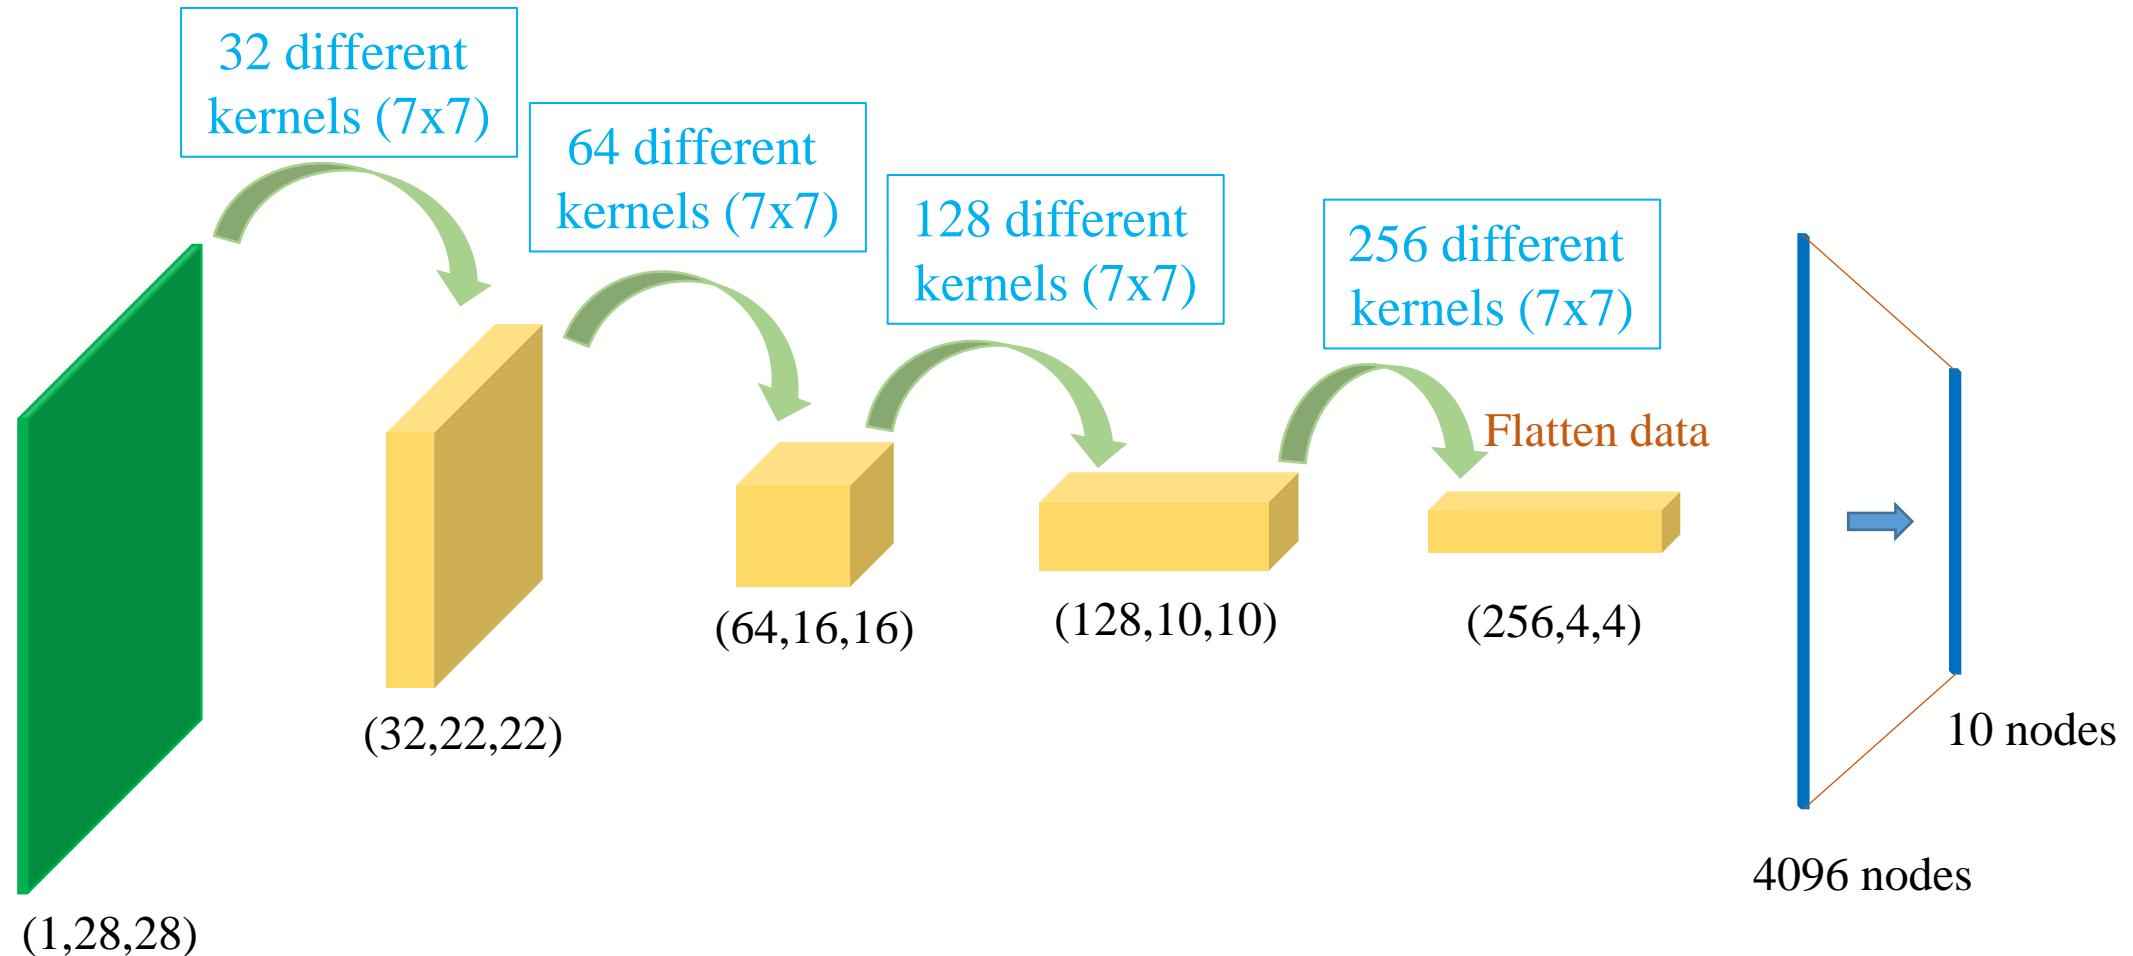
- ❖ Design a model for Fashion-MNIST dataset



Convolutional Neural Network

❖ Design a model for Fashion-MNIST dataset

demo



Simple Convolutional Neural Network

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=7)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=7)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=7)
        self.flatten = nn.Flatten()
        self.dense = nn.Linear(4*4*256, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.relu(self.conv4(x))
        x = self.flatten(x)
        x = self.dense(x)
        return x

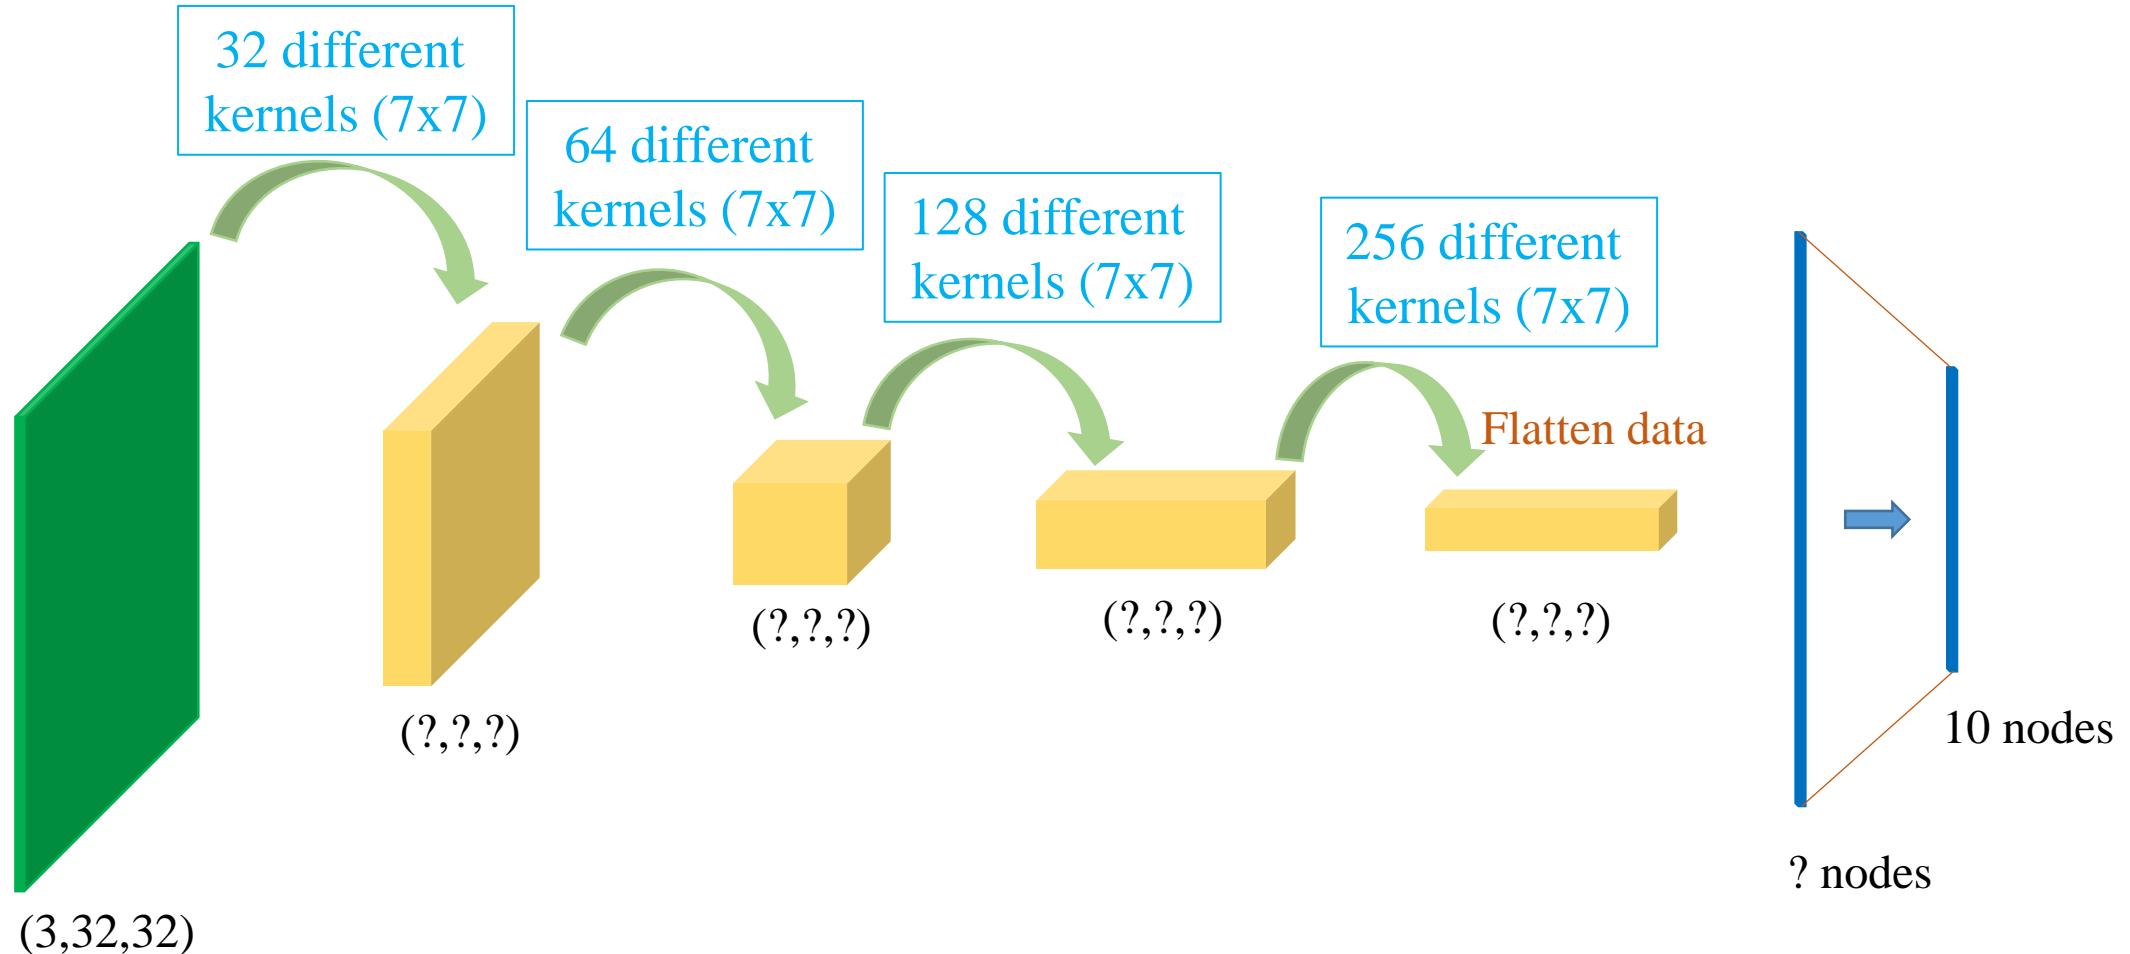
model = CustomModel()
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 22, 22]	1,600
ReLU-2	[-1, 32, 22, 22]	0
Conv2d-3	[-1, 64, 16, 16]	100,416
ReLU-4	[-1, 64, 16, 16]	0
Conv2d-5	[-1, 128, 10, 10]	401,536
ReLU-6	[-1, 128, 10, 10]	0
Conv2d-7	[-1, 256, 4, 4]	1,605,888
ReLU-8	[-1, 256, 4, 4]	0
Flatten-9	[-1, 4096]	0
Linear-10	[-1, 128]	524,416
ReLU-11	[-1, 128]	0
Linear-12	[-1, 10]	1,290
<hr/>		
Total params: 2,635,146		
Trainable params: 2,635,146		
Non-trainable params: 0		
<hr/>		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.78		
Params size (MB): 10.05		
Estimated Total Size (MB): 10.83		
<hr/>		

Convolutional Neural Network

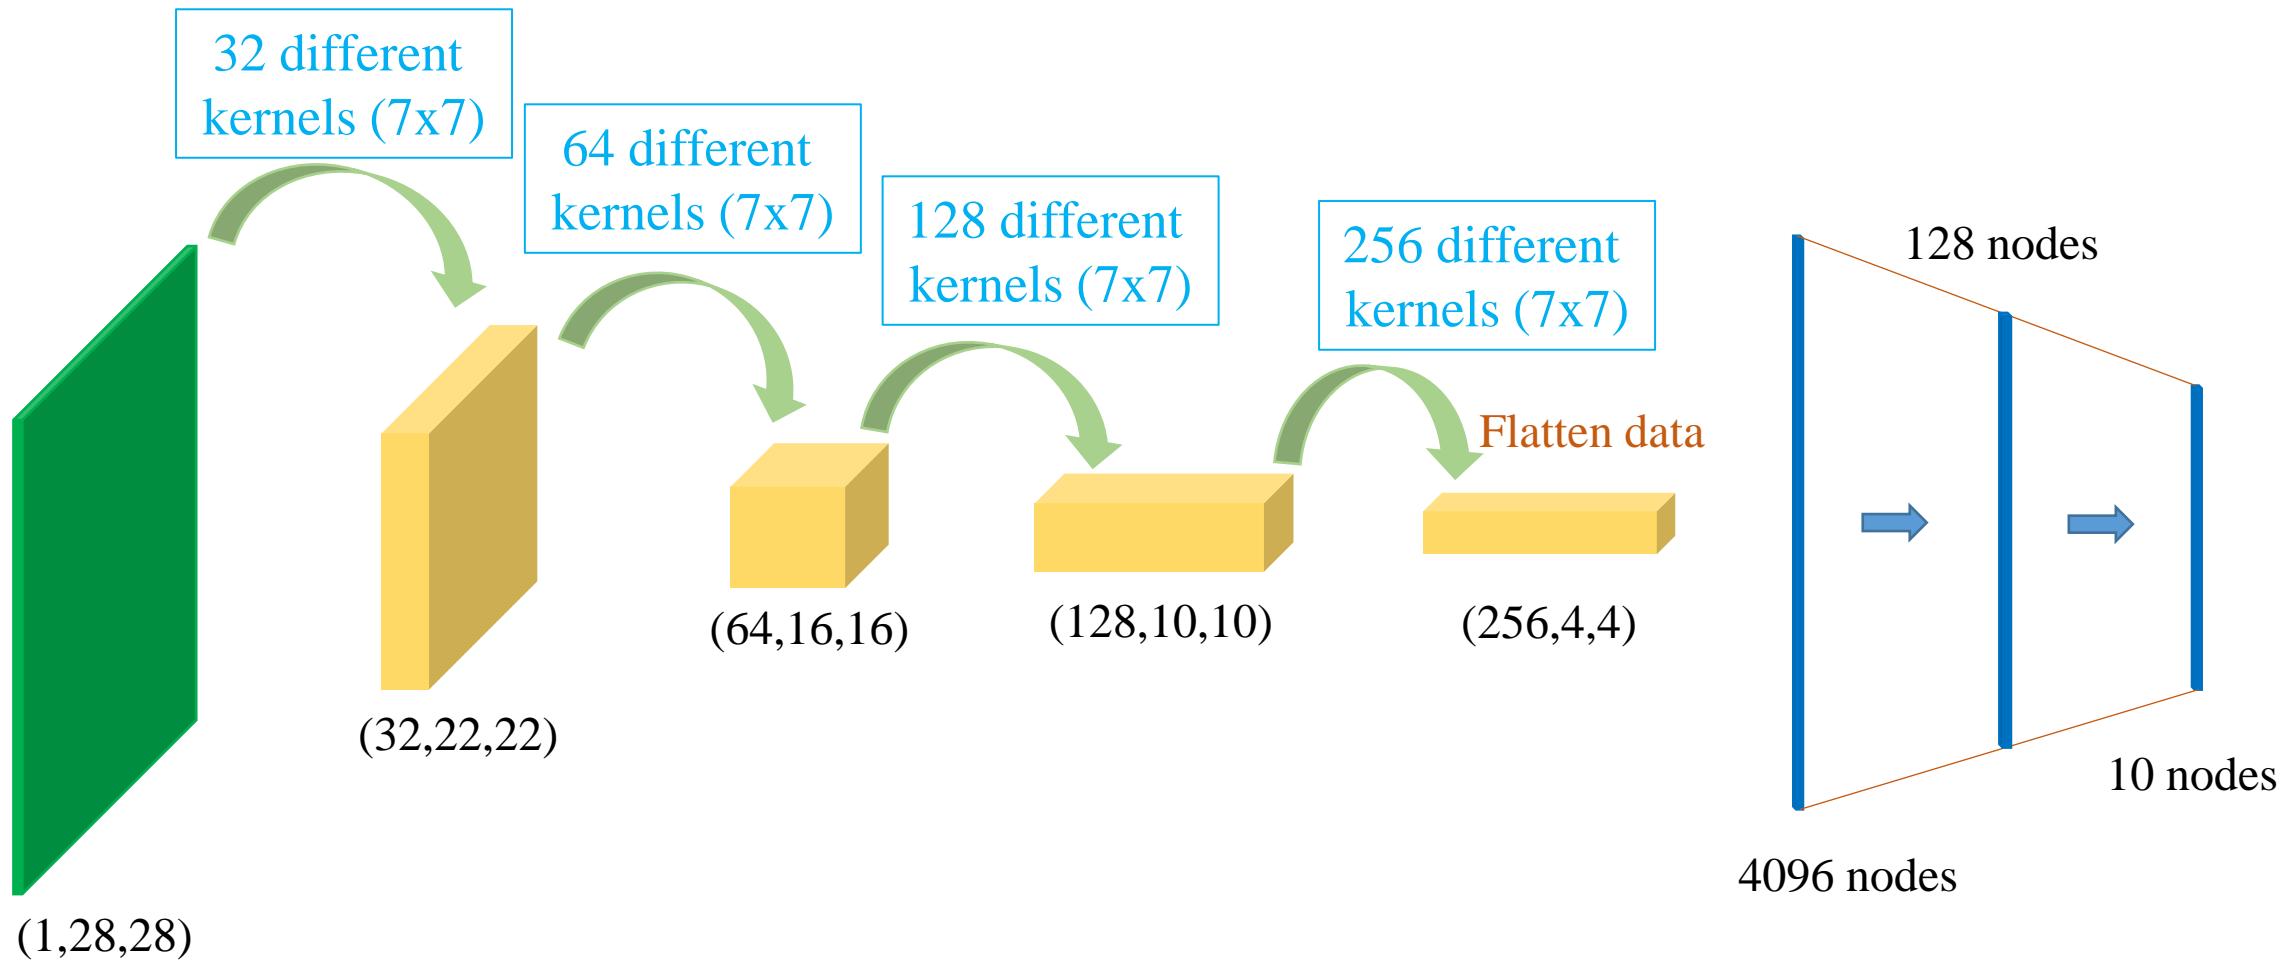
❖ Design a model for Cifar-10 dataset

demo



Convolutional Neural Network

❖ Fashion-MNIST dataset: case 1



```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=7)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=7)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=7)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(4*4*256, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.relu(self.conv4(x))
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x

model = CustomModel()
model = model.to(device)
```

```
# Load FashionMNIST dataset
transform = Compose([ToTensor(),
                     Normalize((0.5,),
                               (0.5,))])

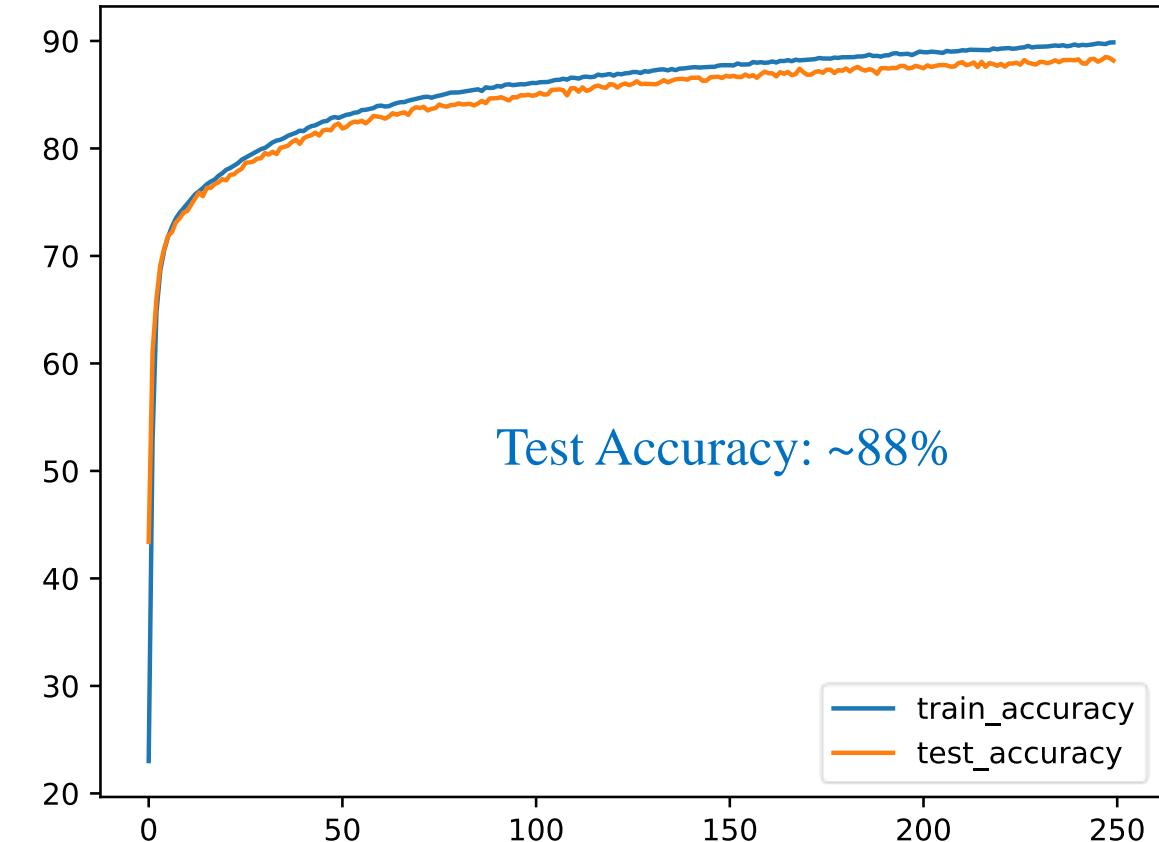
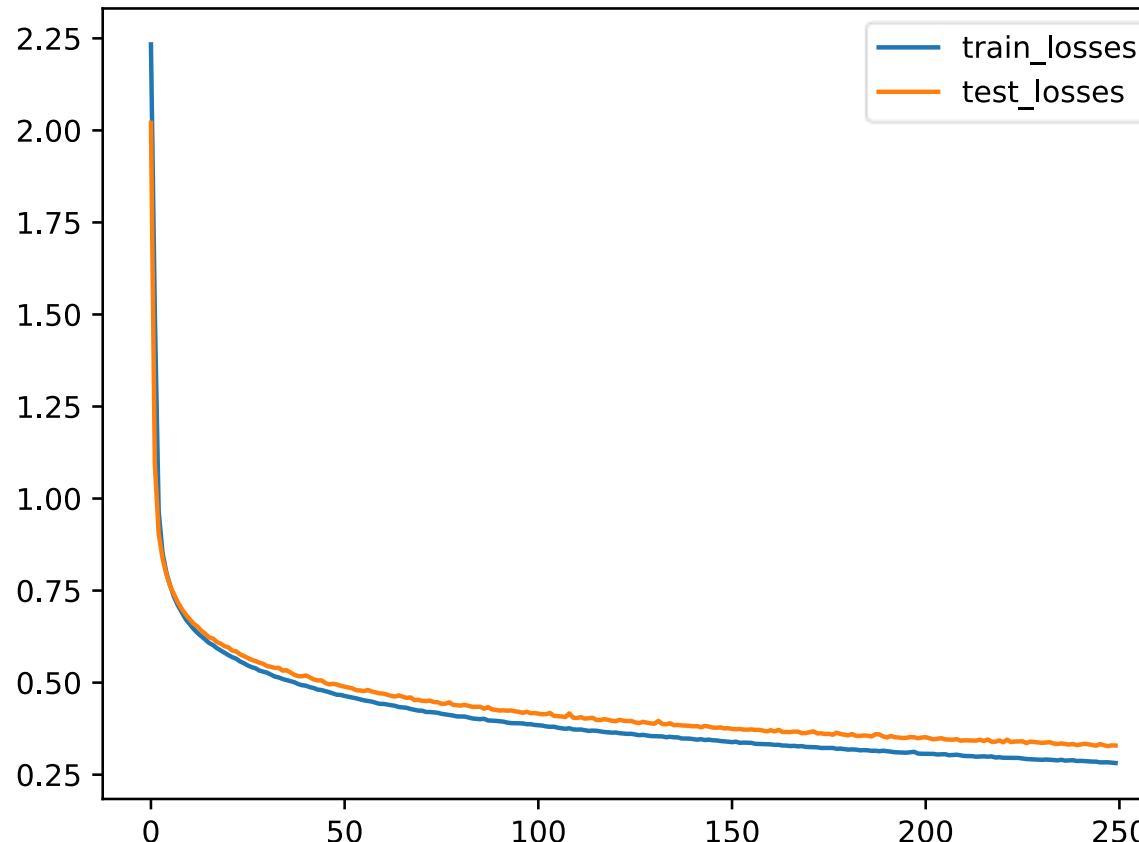
trainset = FashionMNIST(root='data',
                        train=True,
                        download=True,
                        transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = FashionMNIST(root='data',
                       train=False,
                       download=True,
                       transform=transform)
testloader = DataLoader(testset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=False)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-5)
```

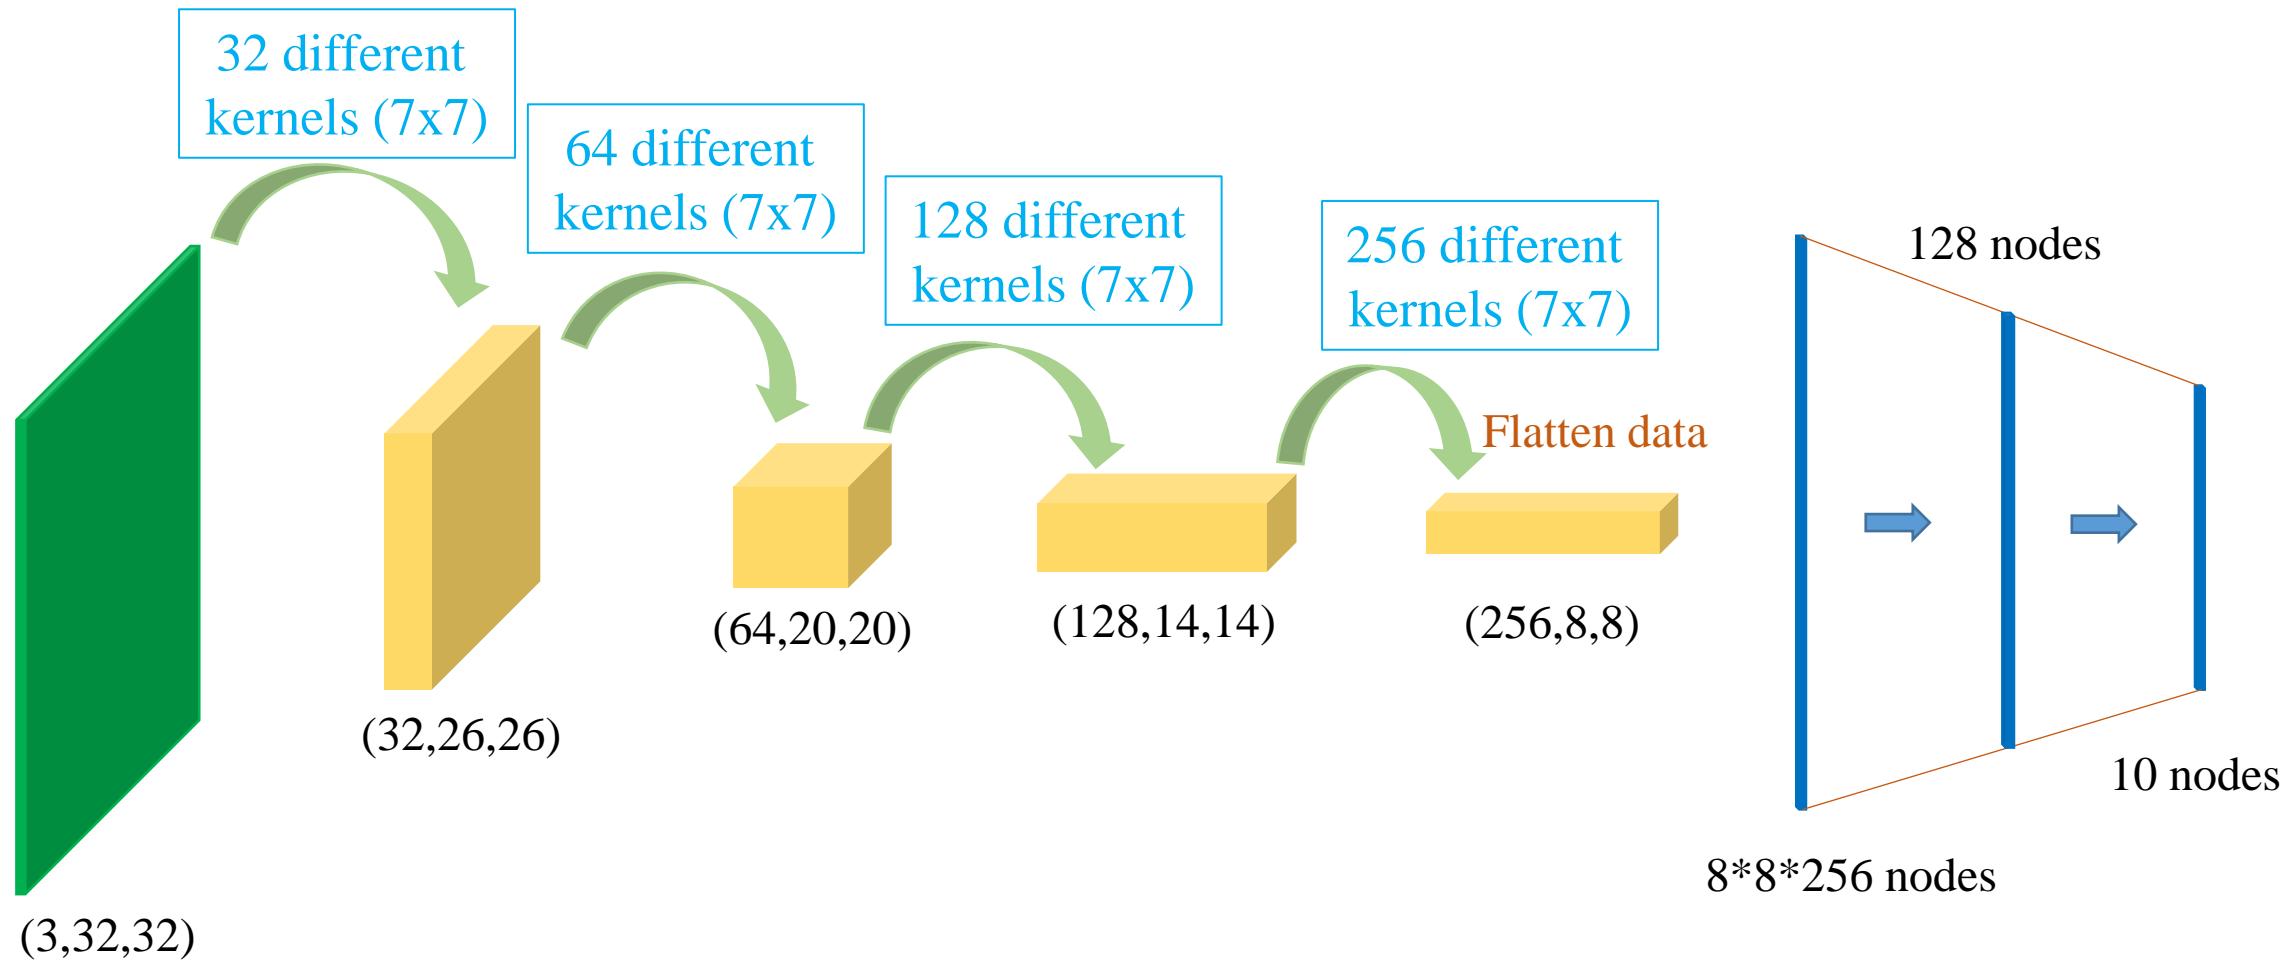
Convolutional Neural Network

❖ Fashion-MNIST dataset: case 1



Convolutional Neural Network

❖ Cifar-10 dataset: case 2



```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=7)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=7)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=7)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(8*8*256, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.relu(self.conv4(x))
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x

model = CustomModel()
model = model.to(device)
```

```
# Load CIFAR10 dataset
transform = Compose([ToTensor(),
                     Normalize((0.5, 0.5, 0.5),
                               (0.5, 0.5, 0.5))])

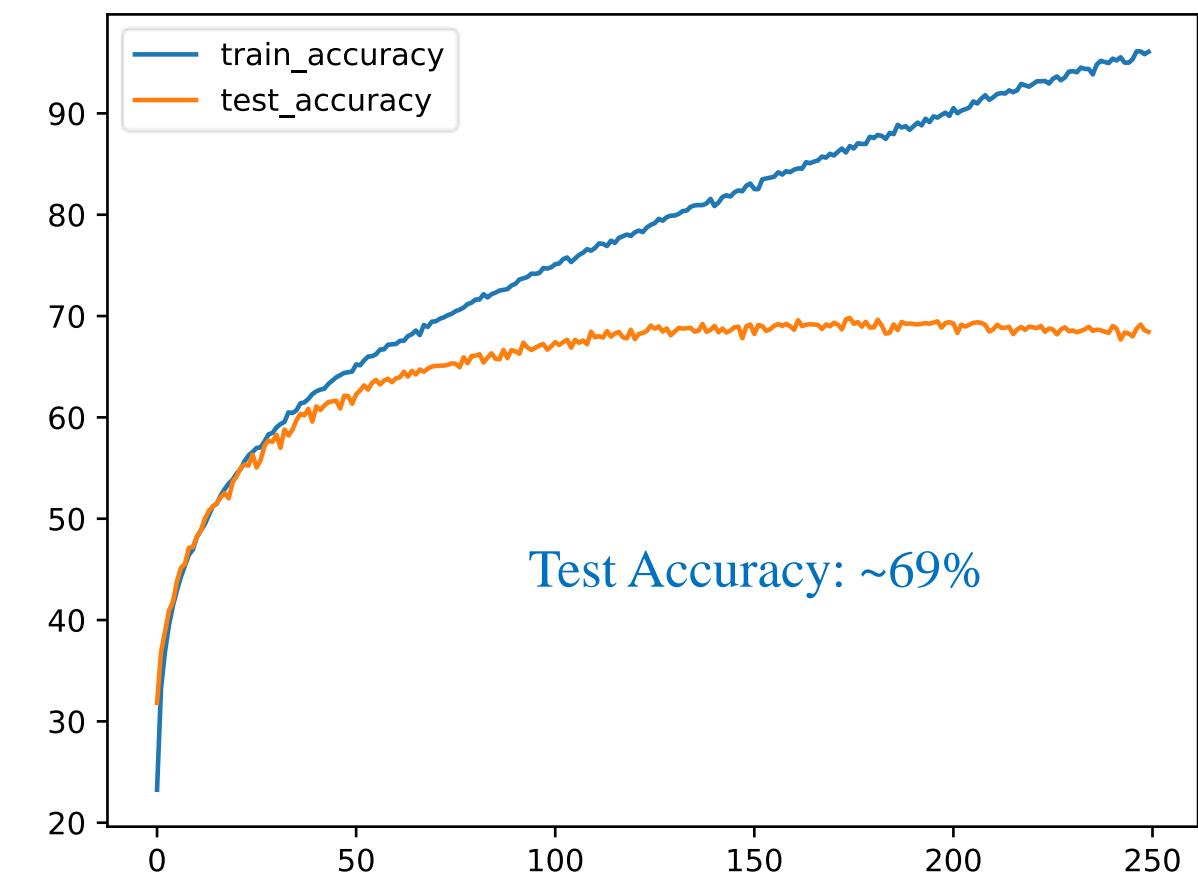
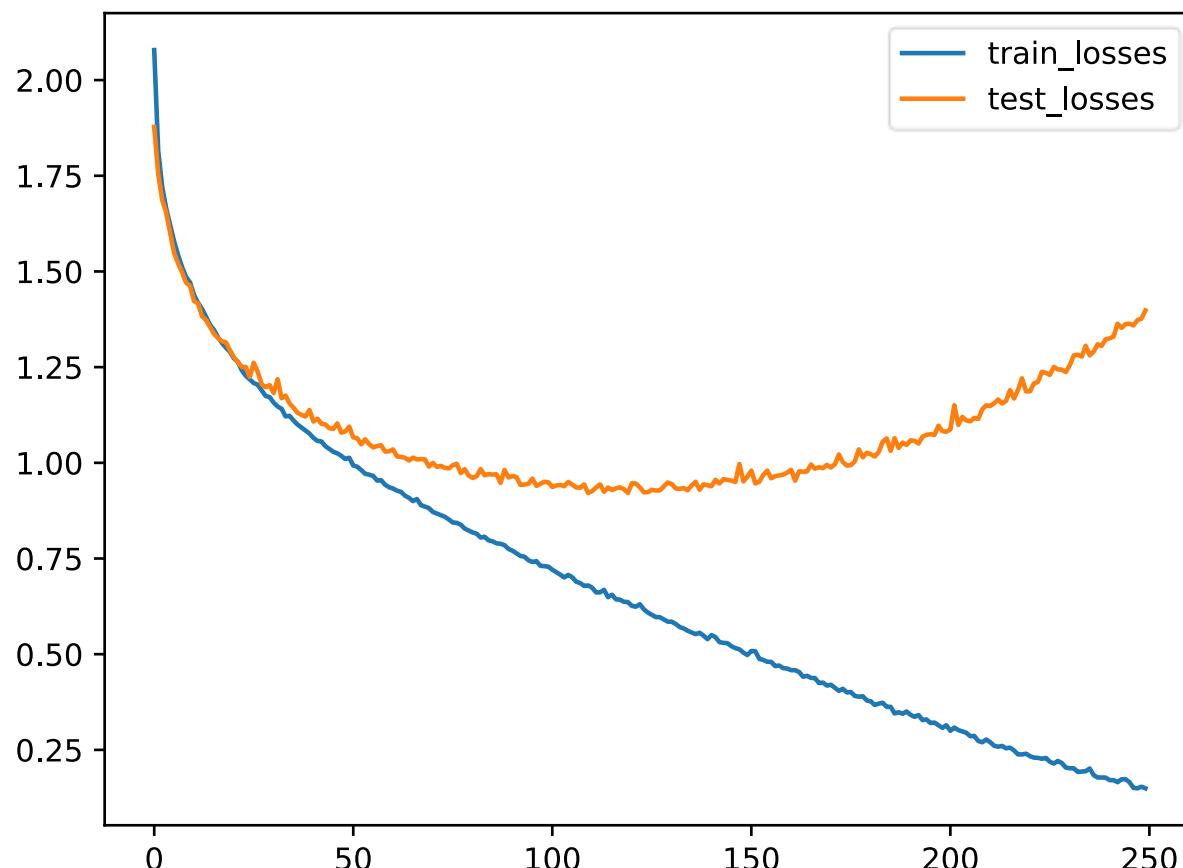
trainset = CIFAR10(root='data',
                   train=True,
                   download=True,
                   transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = CIFAR10(root='data',
                  train=False,
                  download=True,
                  transform=transform)
testloader = DataLoader(testset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=False)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-5)
```

Convolutional Neural Network

❖ Cifar-10 dataset: case 2



Test Accuracy from MLP: ~53%

Outline

- Fashion-MNIST/Cifar10 Using only Conv2d
- Pooling
- Padding
- 1x1 Convolution
- LeNet and VGG Models

Down-sample Feature Map

Max pooling: Features are preserved

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

Data

2x2 max pooling

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

)

=

m_1	m_2
m_3	m_4

$$m_1 = \max(v_1, v_2, v_5, v_6)$$
$$m_2 = \max(v_3, v_4, v_7, v_8)$$
$$m_3 = \max(v_9, v_{10}, v_{13}, v_{14})$$
$$m_4 = \max(v_{11}, v_{12}, v_{15}, v_{16})$$



max pooling
(2x2)



max pooling
(2x2)



`nn.MaxPool2d(2, 2)`

Down-sample Feature Map

Average pooling: Features are preserved

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

Data

average pooling (

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

) =

m_1	m_2
m_3	m_4

$$m_1 = \text{mean}(v_1, v_2, v_5, v_6)$$
$$m_2 = \text{mean}(v_3, v_4, v_7, v_8)$$
$$m_3 = \text{mean}(v_9, v_{10}, v_{13}, v_{14})$$
$$m_4 = \text{mean}(v_{11}, v_{12}, v_{15}, v_{16})$$

`nn.AvgPool2d(2, 2)`



Average
Pooling (2x2) \rightarrow



Average
Pooling (2x2) \rightarrow



Down-sample Feature Map

Max pooling vs. Average pooling

`nn.MaxPool2d(2, 2)`



Feature map (220x220)

max pooling
(2x2)



max pooling
(2x2)



Feature map
(55x55)

`nn.AvgPool2d(2, 2)`



Feature map (220x220)

Average
Pooling (2x2)



Average
Pooling (2x2)



Feature map
(55x55)

```

def forward(self, x):
    x = self.relu(self.conv1(x))
    x = self.relu(self.conv2(x))
    x = self.pool(x)

    x = self.relu(self.conv3(x))
    x = self.relu(self.conv4(x))
    x = self.pool(x)

    x = self.relu(self.conv5(x))
    x = self.flatten(x)
    x = self.dense(x)

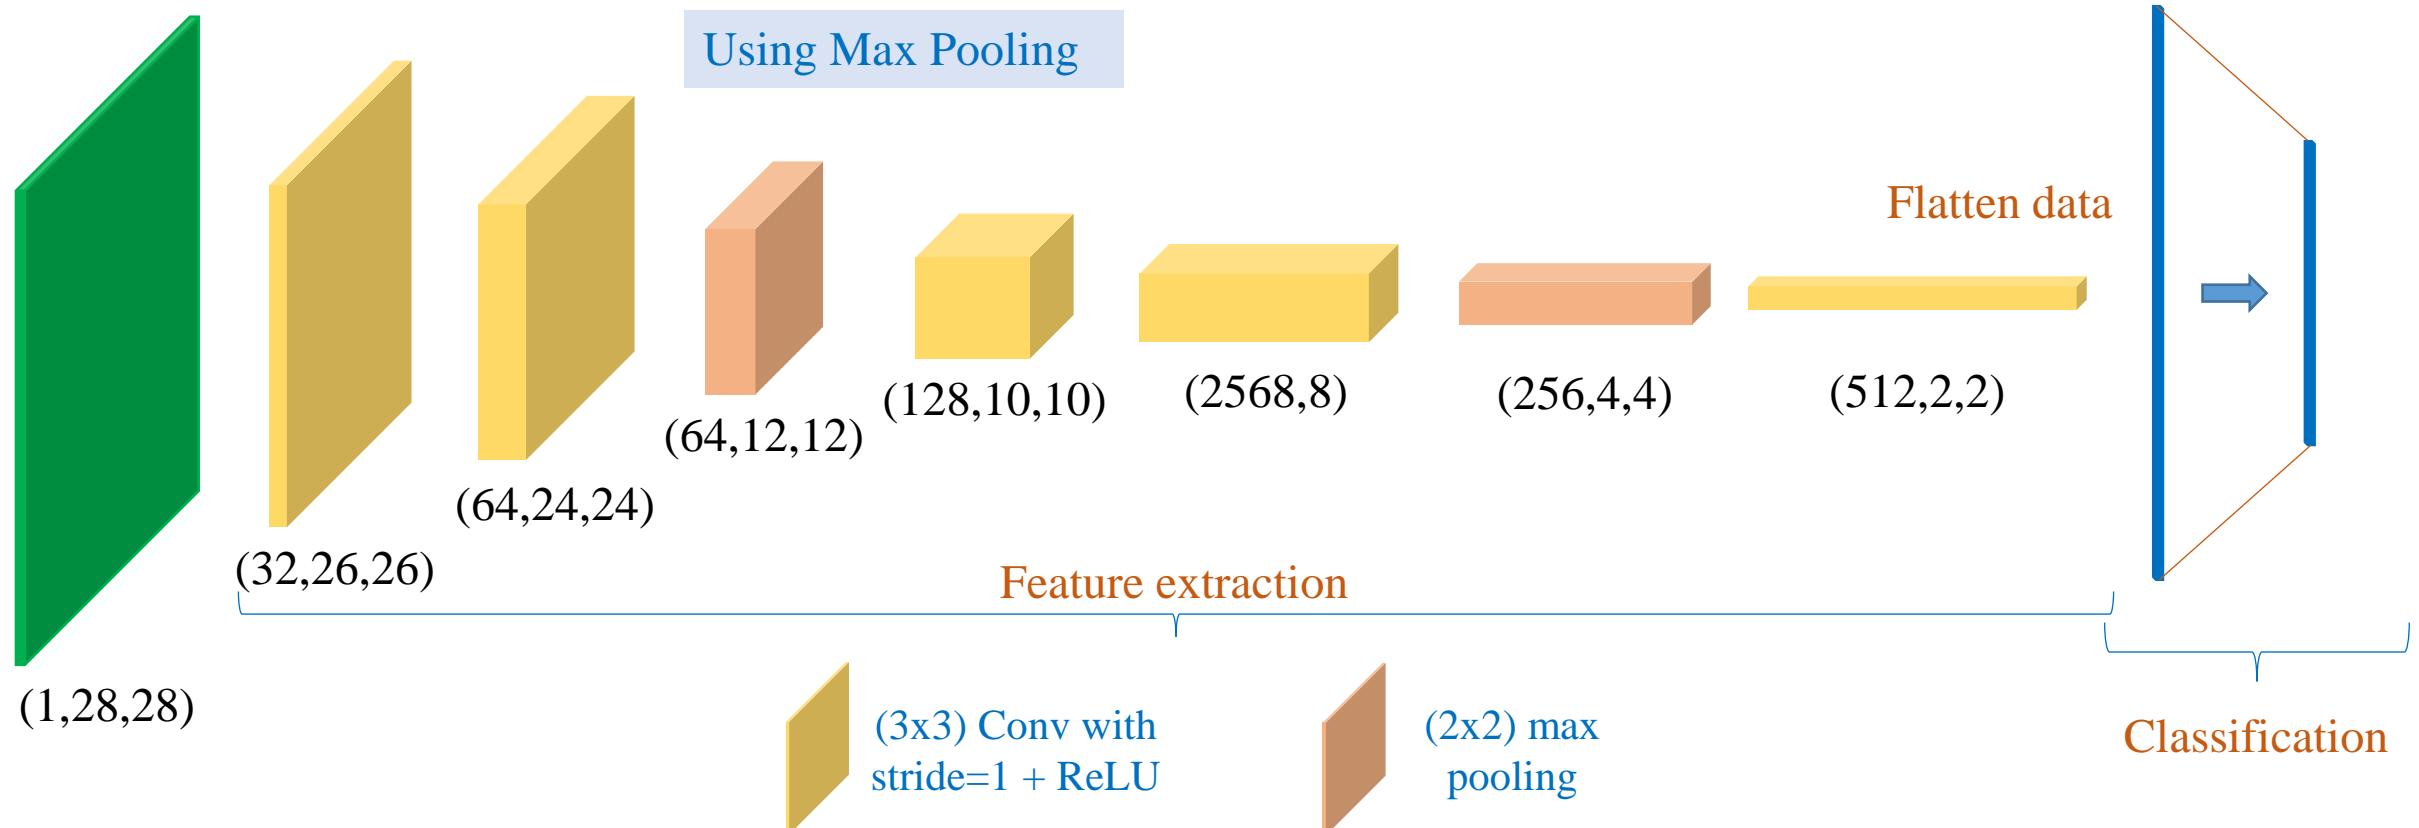
    return x

```

```

class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=3)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=3)
        self.conv5 = nn.Conv2d(256, 512, kernel_size=3)
        self.relu = nn.ReLU()
        self.pool = nn.MaxPool2d(2, 2) # 2x2 Max pooling
        self.flatten = nn.Flatten()
        self.dense = nn.Linear(2*2*512, 10)

```



Down-sample Feature Map

❖ Convolve with stride

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

Data D

w_1	w_2
w_3	w_4

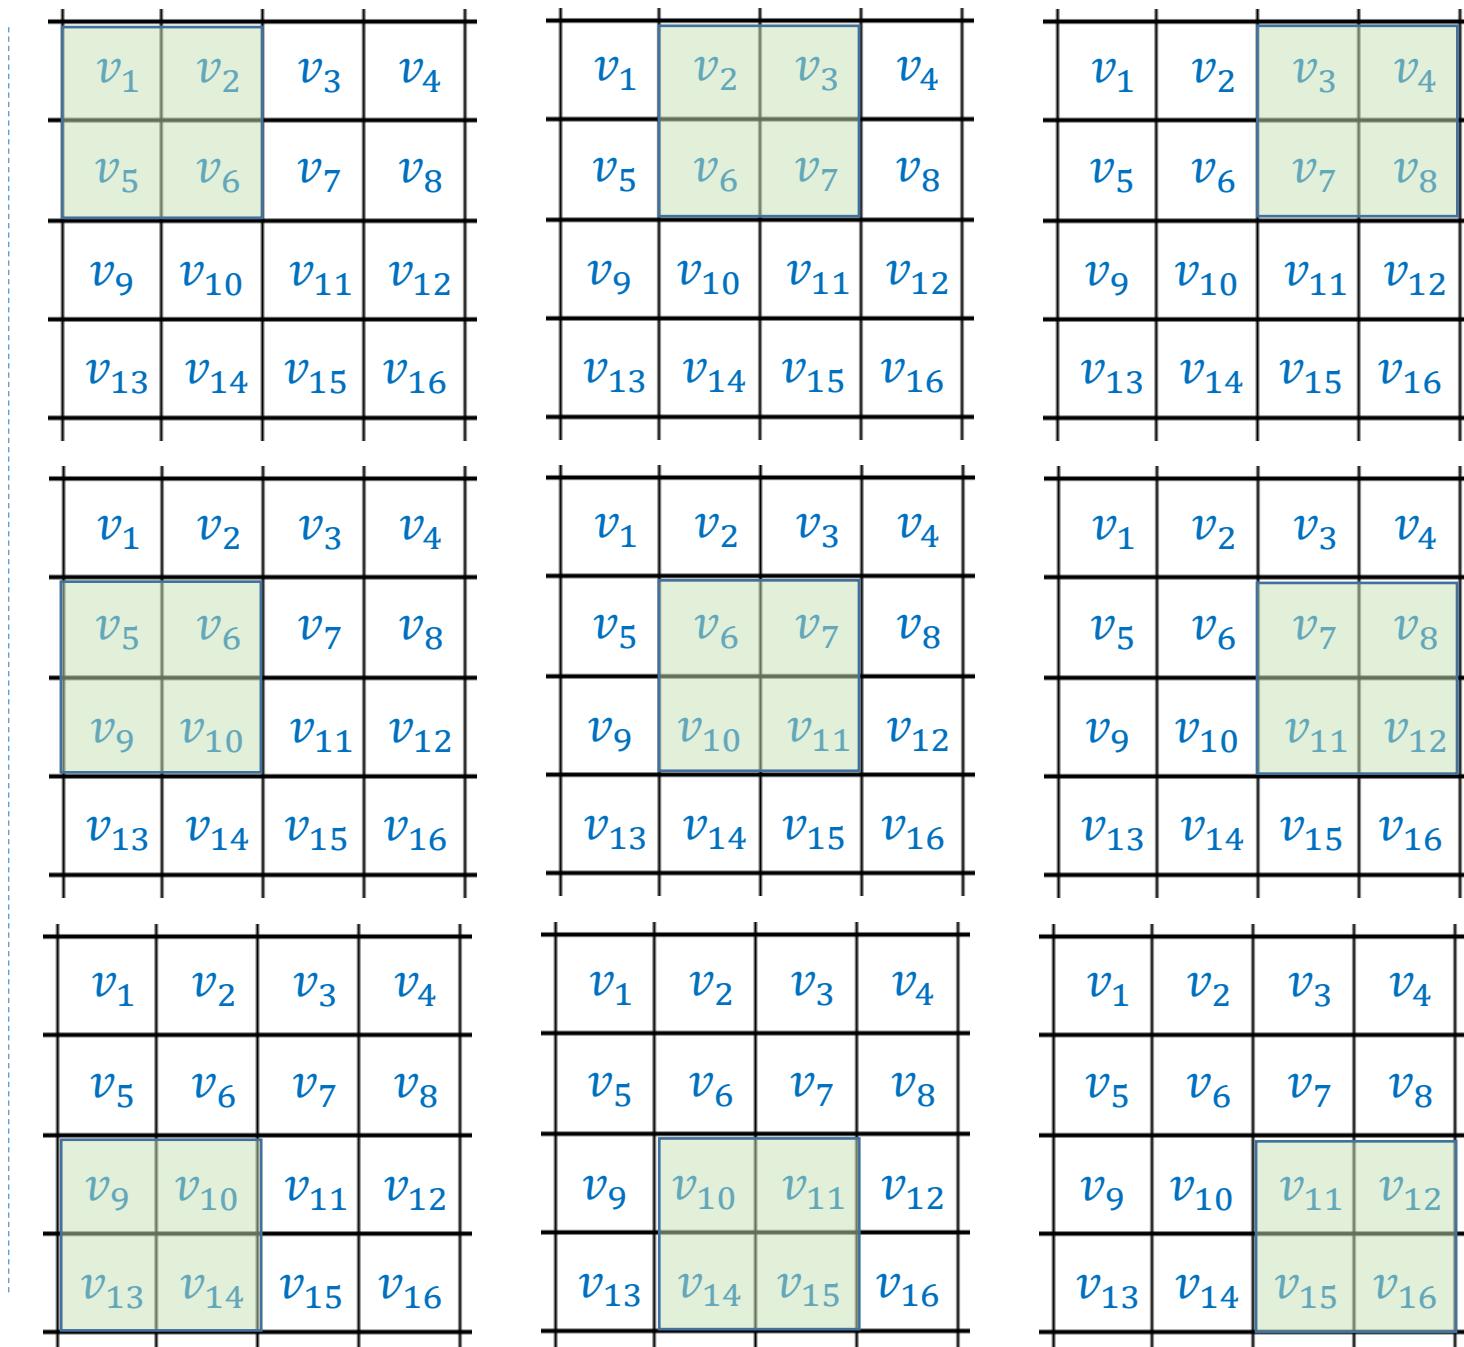
b

Kernel of parameters

Convolve D
with stride=1

m_1	m_2	m_3
m_4	m_5	m_6
m_7	m_8	m_9

Output



Down-sample Feature Map

❖ Convolve with stride

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

Data D

w_1	w_2
w_3	w_4

b

Kernel of parameters

Convolve D
with stride=2

m_1	m_2
m_3	m_4

Output

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

`nn.Conv2d(in_channels, out_channels, kernel_size, stride=1)`

`nn.Conv2d(in_channels, out_channels, kernel_size, stride=2)`

```

class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=3)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=2)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=3, stride=2)
        self.conv5 = nn.Conv2d(256, 512, kernel_size=3)
        self.relu = nn.ReLU()
        self.flatten = nn.Flatten()
        self.dense = nn.Linear(2*2*512, 10)

```

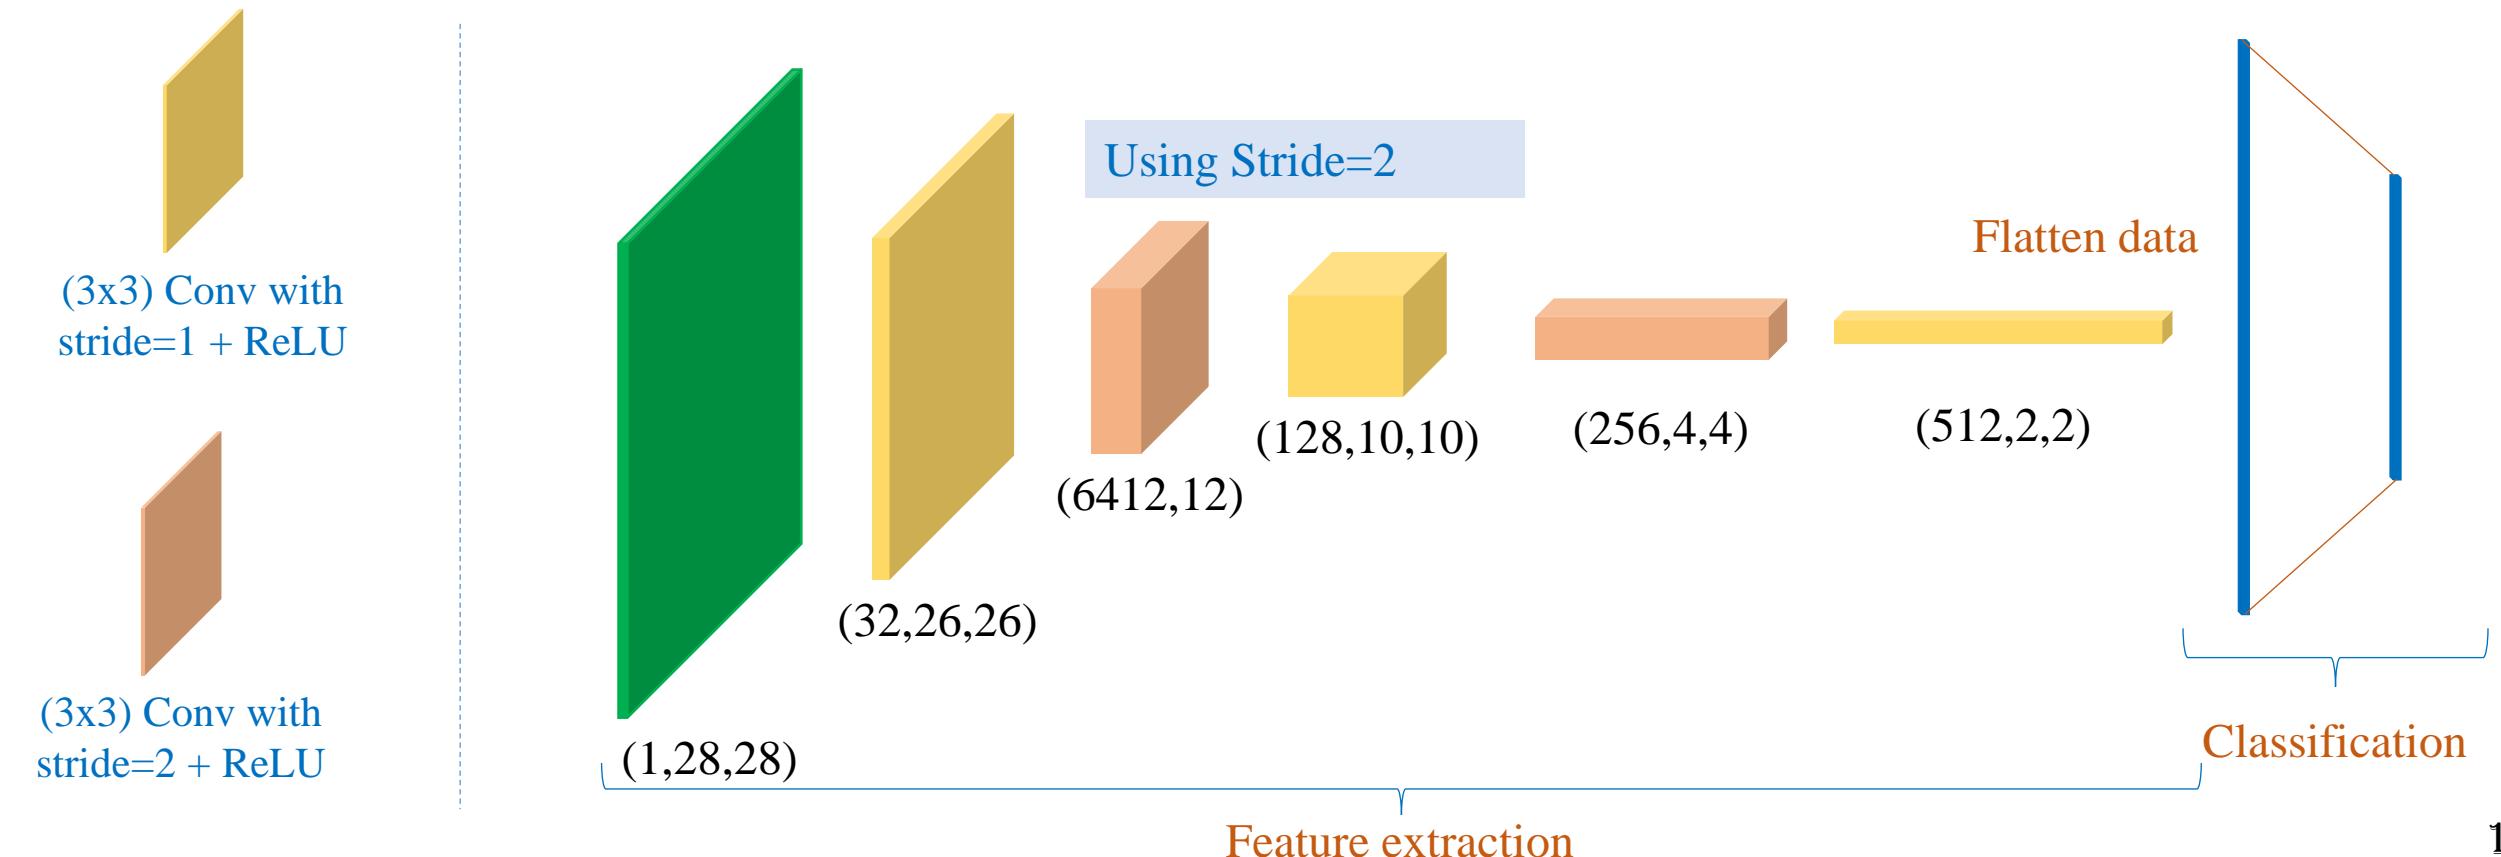
```

def forward(self, x):
    x = self.relu(self.conv1(x))
    x = self.relu(self.conv2(x))

    x = self.relu(self.conv3(x))
    x = self.relu(self.conv4(x))

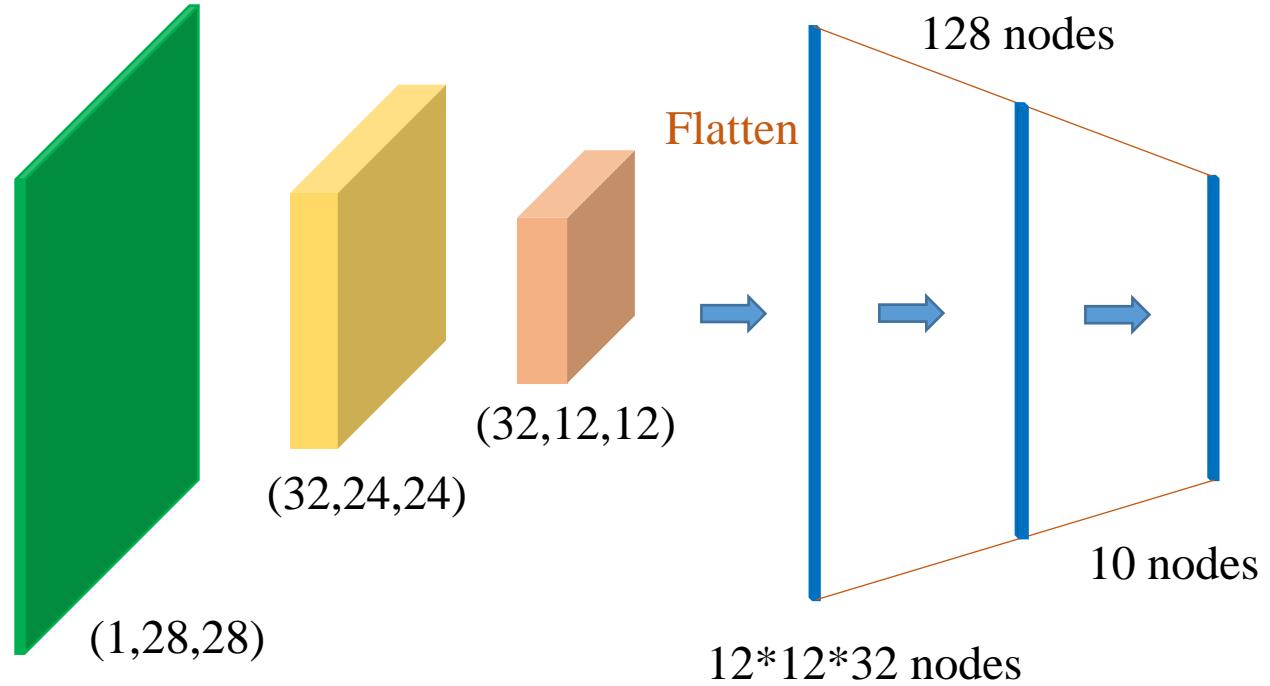
    x = self.relu(self.conv5(x))
    x = self.flatten(x)
    x = self.dense(x)
    return x

```



Comparison of down-samplings

Model Design

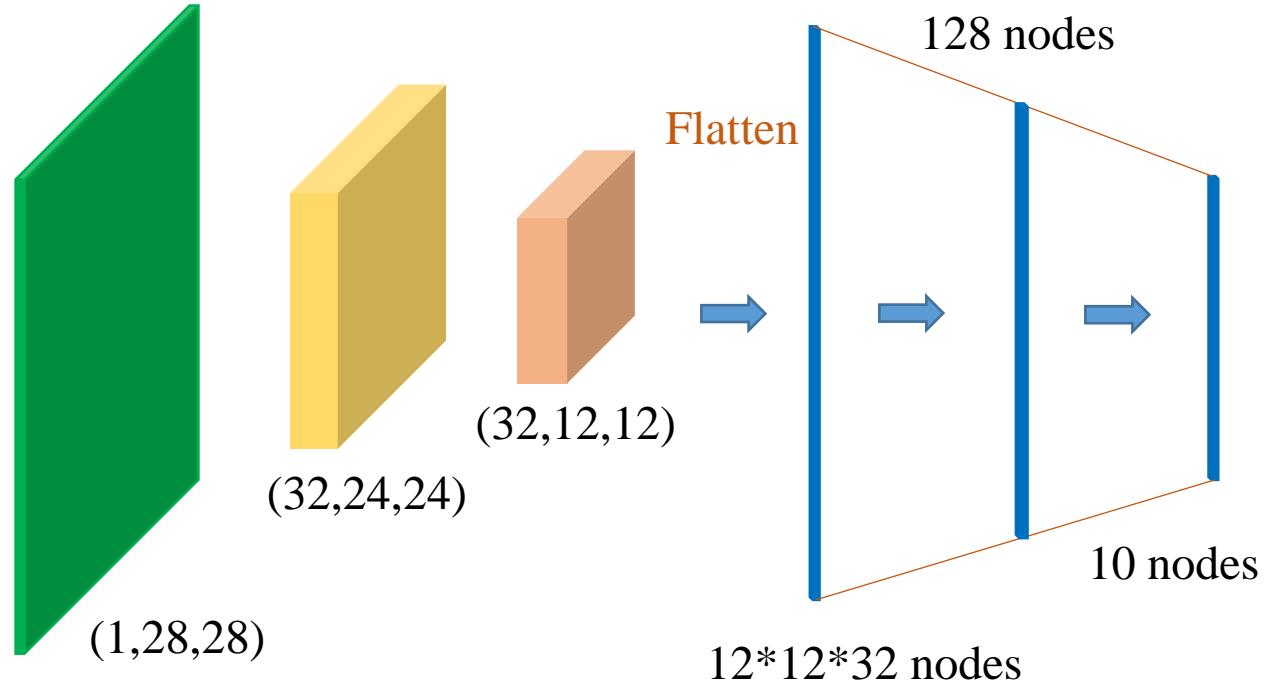


Implementation 1

```
class CustomModel(nn.Module):  
    def __init__(self):  
        super(CustomModel, self).__init__()  
        self.conv = nn.Conv2d(1, 32, kernel_size=5)  
        self.pool = nn.MaxPool2d(2, 2)  
        self.flatten = nn.Flatten()  
        self.dense1 = nn.Linear(12*12*32, 128)  
        self.dense2 = nn.Linear(128, 10)  
        self.relu = nn.ReLU()  
  
    def forward(self, x):  
        x = self.conv(x)  
        x = self.relu(x)  
        x = self.pool(x)  
        x = self.flatten(x)  
        x = self.relu(self.dense1(x))  
        x = self.dense2(x)  
        return x  
  
# create model  
model = CustomModel()  
model = model.to(device)  
  
# Loss and optimizer  
criterion = nn.CrossEntropyLoss()  
optimizer = optim.Adam(model.parameters(), lr=1e-5)
```

Comparison of down-samplings

Model Design



(5x5) Conv with
stride=1 + ReLU

(2x2) down-
sampling

Implementation 2

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv = nn.Conv2d(1, 32, kernel_size=5)
        self.pool = nn.AvgPool2d(2, 2)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(12*12*32, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.conv(x)
        x = self.relu(x)
        x = self.pool(x)
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x

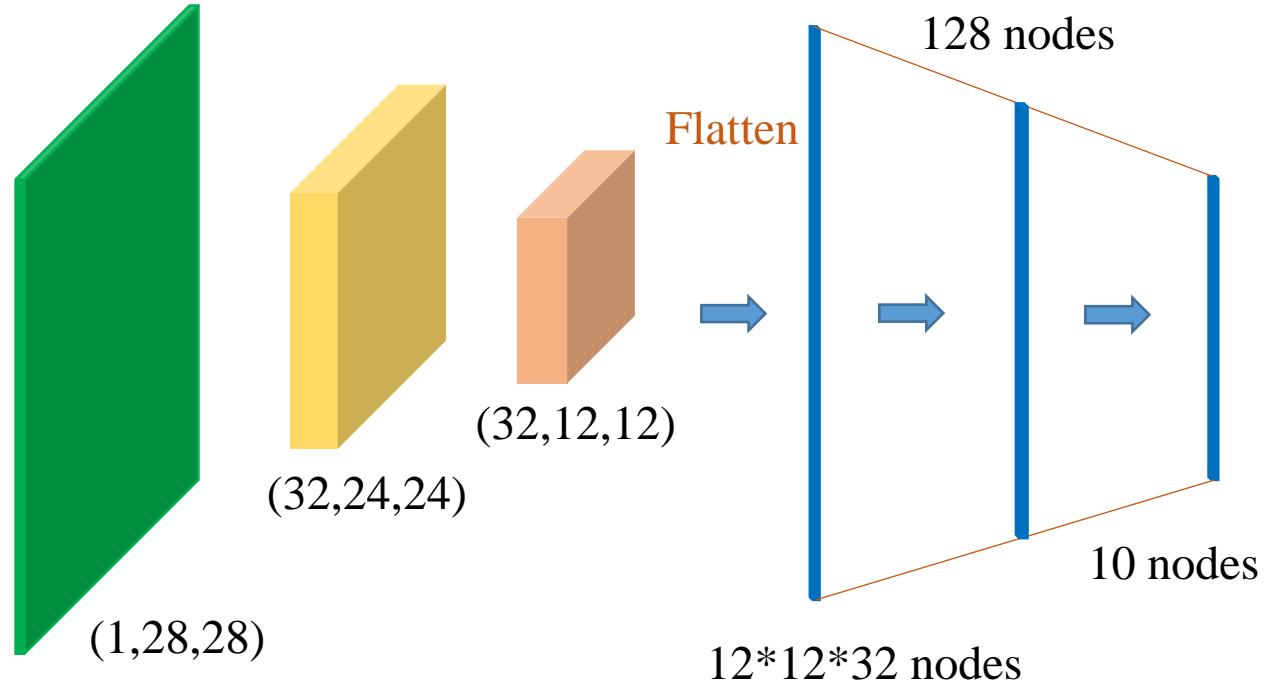
# create model
model = CustomModel()
model = model.to(device)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-5)
```

average pooling

Comparison of down-samplings

Model Design



(5x5) Conv with
stride=1 + ReLU

(2x2) down-
sampling

```
class ResizeLayer(nn.Module):
    def __init__(self, scale_factor, mode='bilinear',
                 align_corners=False):
        super(ResizeLayer, self).__init__()
        self.scale_factor = scale_factor
        self.mode = mode
        self.align_corners = align_corners

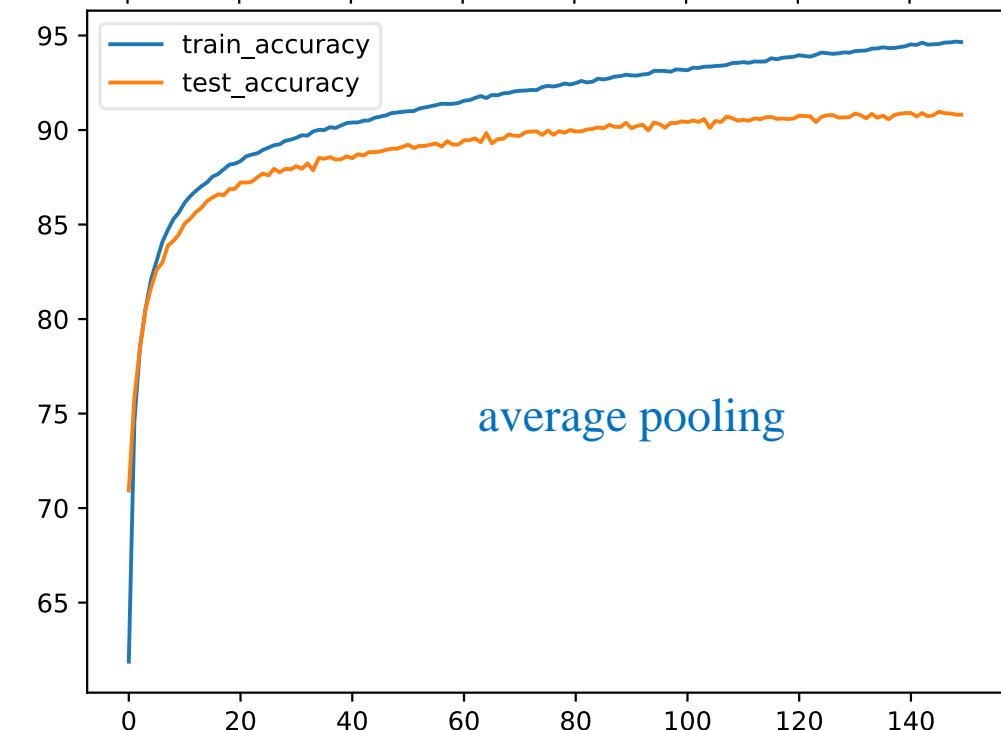
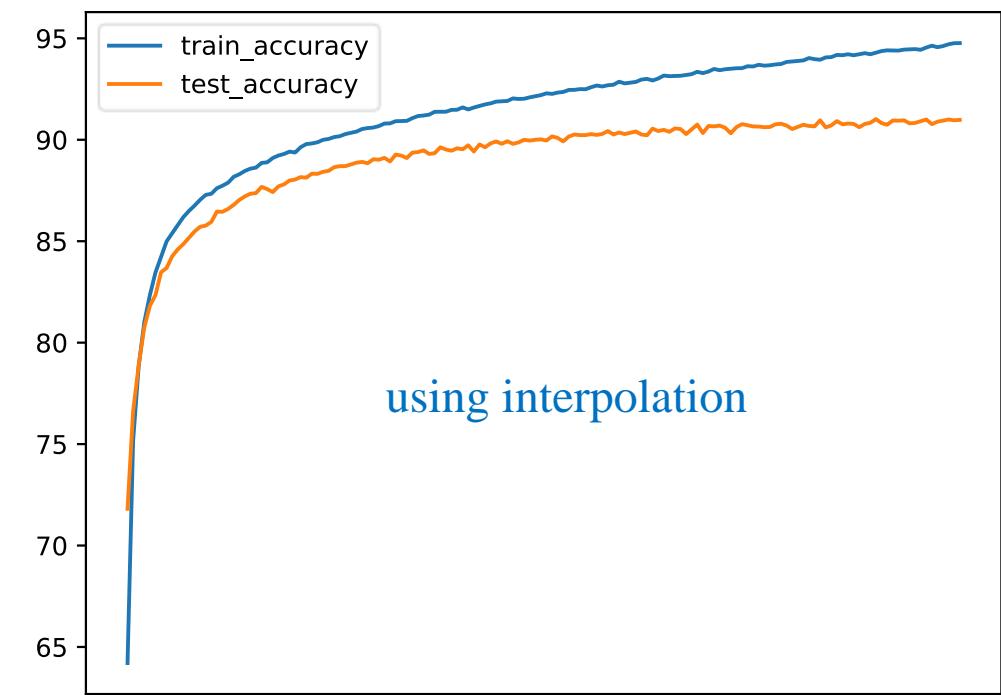
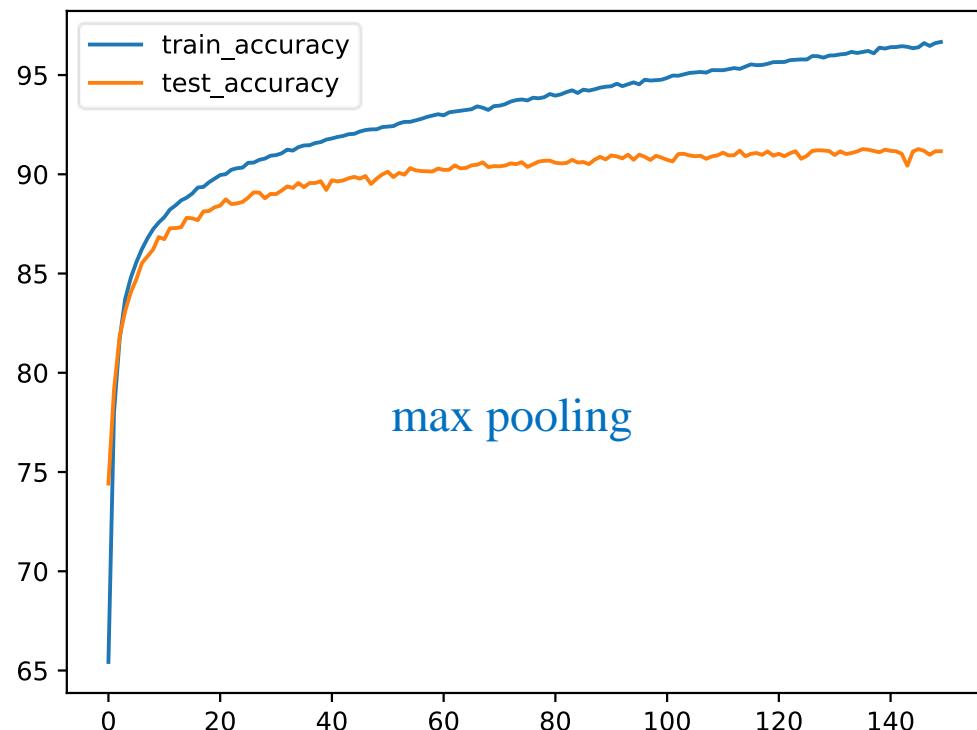
    def forward(self, x):
        return F.interpolate(x, scale_factor=self.scale_factor,
                           mode=self.mode,
                           align_corners=self.align_corners)
```

interpolation

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv = nn.Conv2d(1, 32, kernel_size=5)
        self.resize = ResizeLayer(0.5)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(12*12*32, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv(x))
        x = self.resize(x)
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x
```

Comparison of down-samplings



Padding

$$S_o = \left\lfloor \frac{S_D - K + 2P}{S} \right\rfloor + 1$$

Keep resolution
of feature map

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

Data D
(4x4)

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

b

Kernel of parameters

Without using padding
or padding=0

Padding = 1

v	v	v	v	v	v
v	v_1	v_2	v_3	v_4	v
v	v_5	v_6	v_7	v_8	v
v	v_9	v_{10}	v_{11}	v_{12}	v
v	v_{13}	v_{14}	v_{15}	v_{16}	v
v	v	v	v	v	v

Data D_p

Convolve
with stride=1 (D) =

m_1	m_2
m_4	m_5

Output
(2x2)

Convolve
with stride=1 (D_p) =

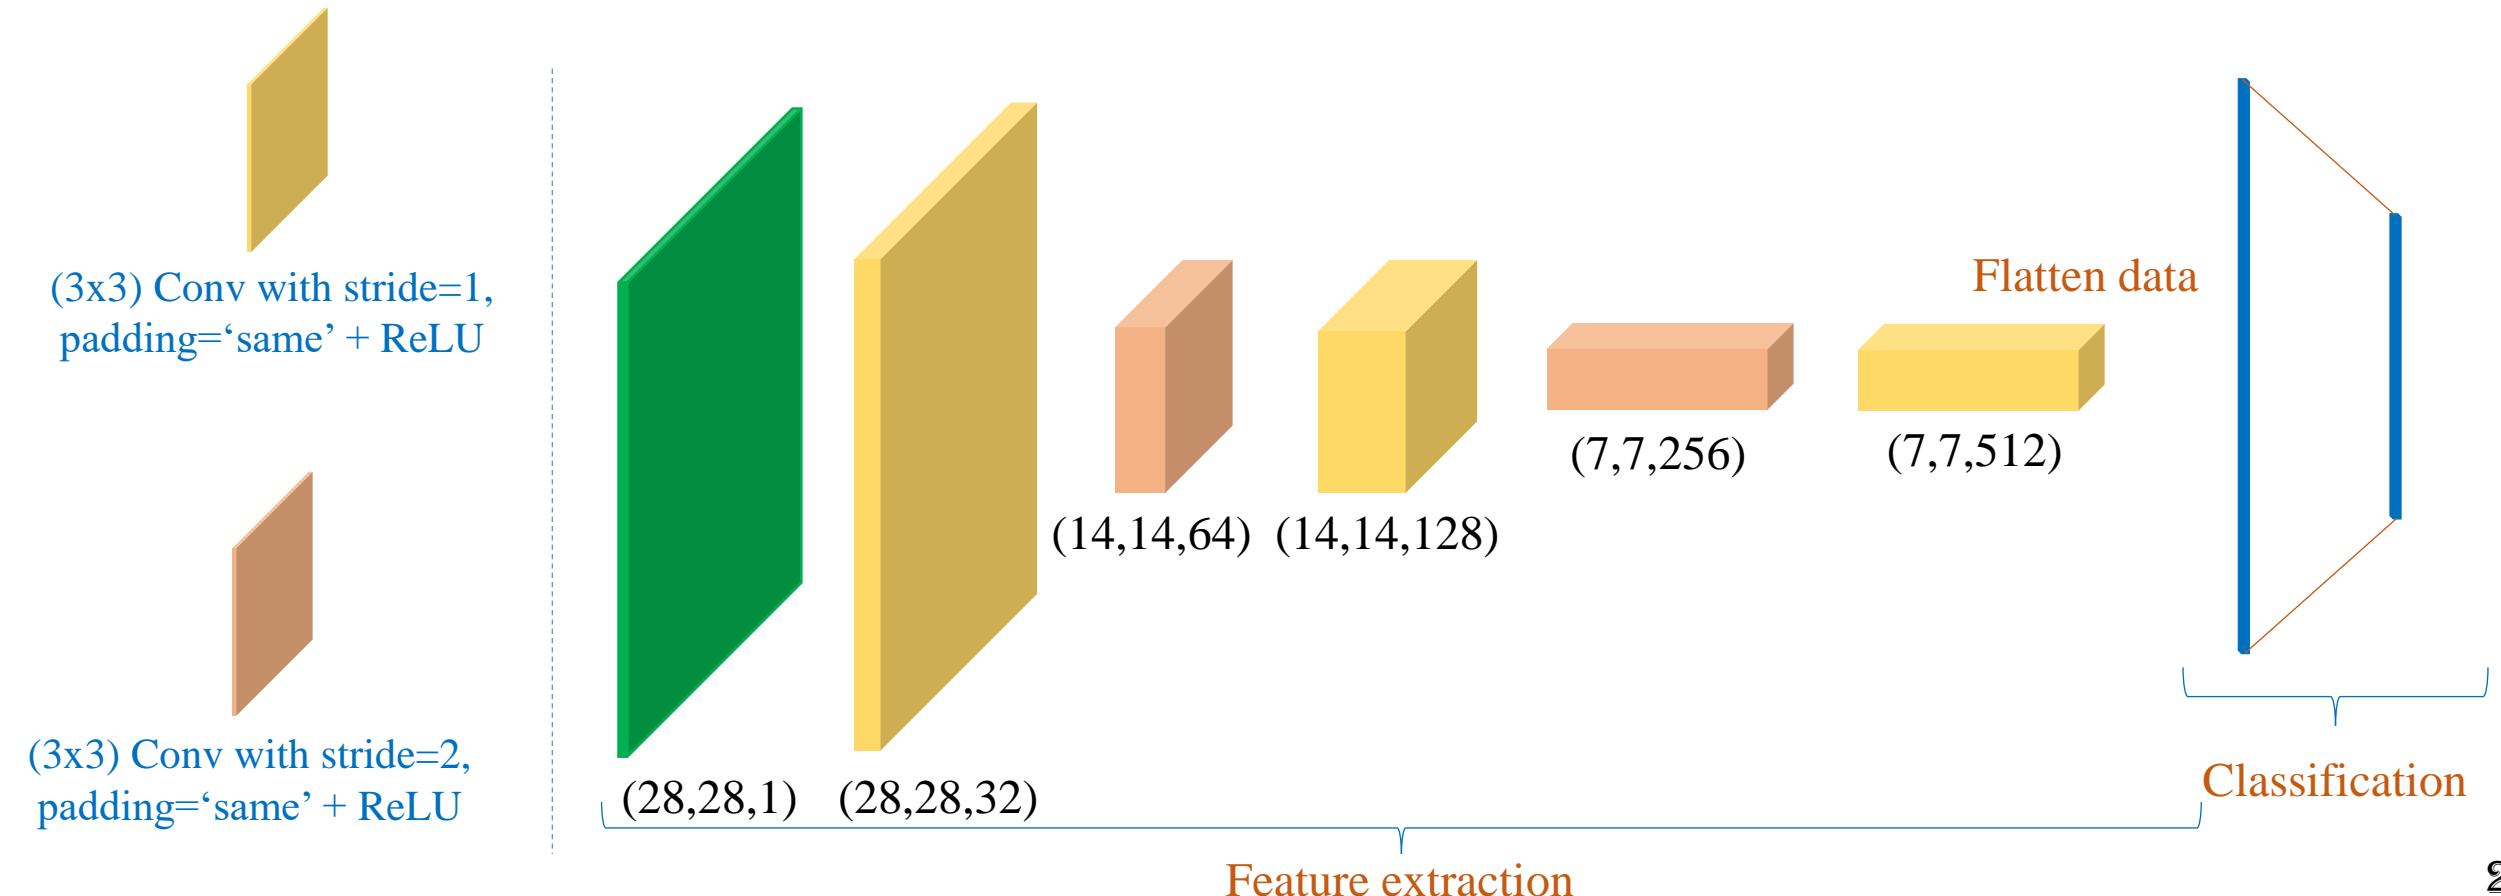
m_1	m_2	m_3	m_4
m_5	m_6	m_7	m_8
m_9	m_{10}	m_{11}	m_{12}
m_{13}	m_{14}	m_{15}	m_{16}

Output
(4x4)

Padding

Example

```
model = nn.Sequential(  
    nn.Conv2d(1, 32, kernel_size=3, padding='same', stride=1), nn.ReLU(),  
    nn.Conv2d(32, 64, kernel_size=3, padding=1, stride=2), nn.ReLU(),  
    nn.Conv2d(64, 128, kernel_size=3, padding=1, stride=1), nn.ReLU(),  
    nn.Conv2d(128, 256, kernel_size=3, padding=1, stride=2), nn.ReLU(),  
    nn.Conv2d(256, 512, kernel_size=3, padding=1, stride=1), nn.ReLU(),  
    nn.Flatten(), nn.Linear(7*7*512, 10)  
)
```



Outline

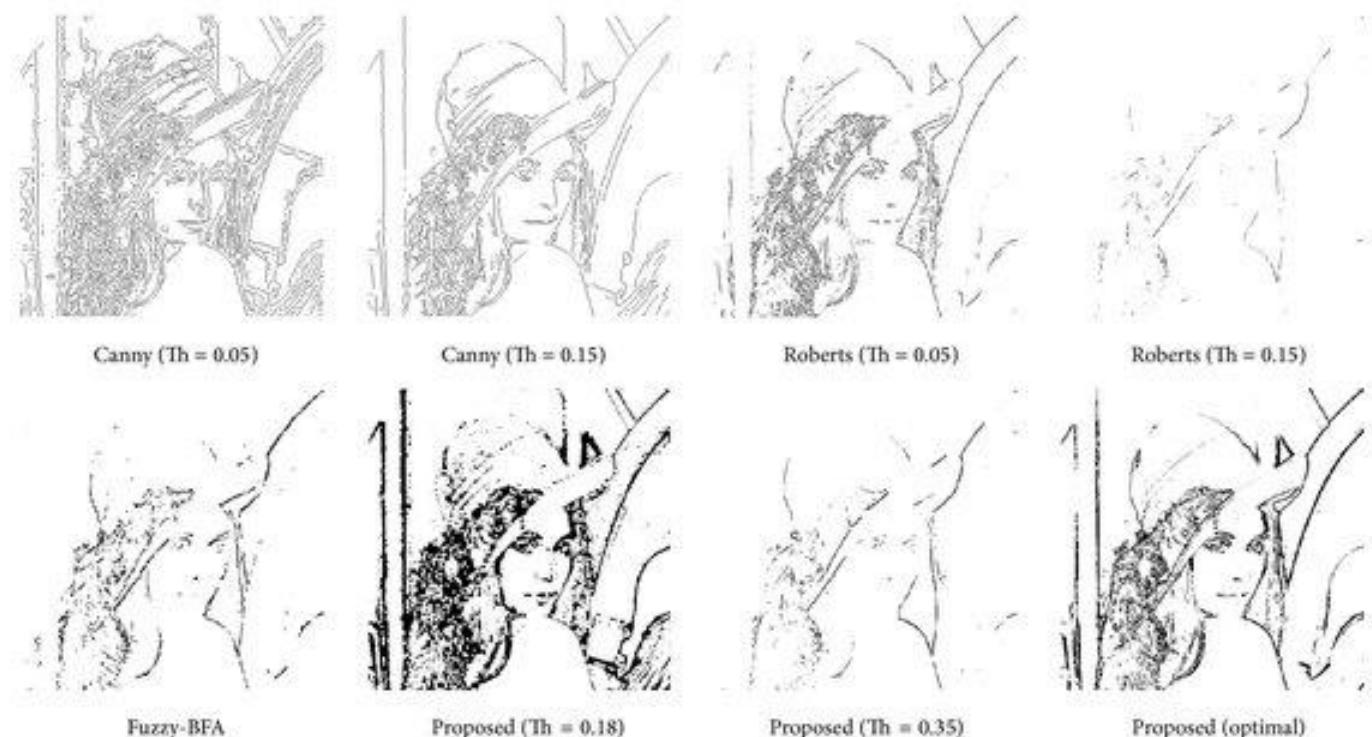
- Fashion-MNIST/Cifar10 Using only Conv2d
- Pooling
- Padding
- 1x1 Convolution
- LeNet and VGG Models

Engineering Feature Extractors

❖ Edge detection



❖ Gradient



Edge Detection via Edge-Strength Estimation
Using Fuzzy Reasoning and Optimal Threshold
Selection Using Particle Swarm Optimization

Derivative and Applications

Tính đạo hàm trung bình theo hướng x

$$\begin{array}{|c|c|c|} \hline 1 & & \\ \hline 2 & & \\ \hline 1 & & \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$

weighted average x-derivative Sobel for x direction

Tính đạo hàm trung bình theo hướng y

$$\begin{array}{|c|c|c|} \hline 1 & & \\ \hline 0 & & \\ \hline -1 & & \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

y-derivative weighted average Sobel for y direction

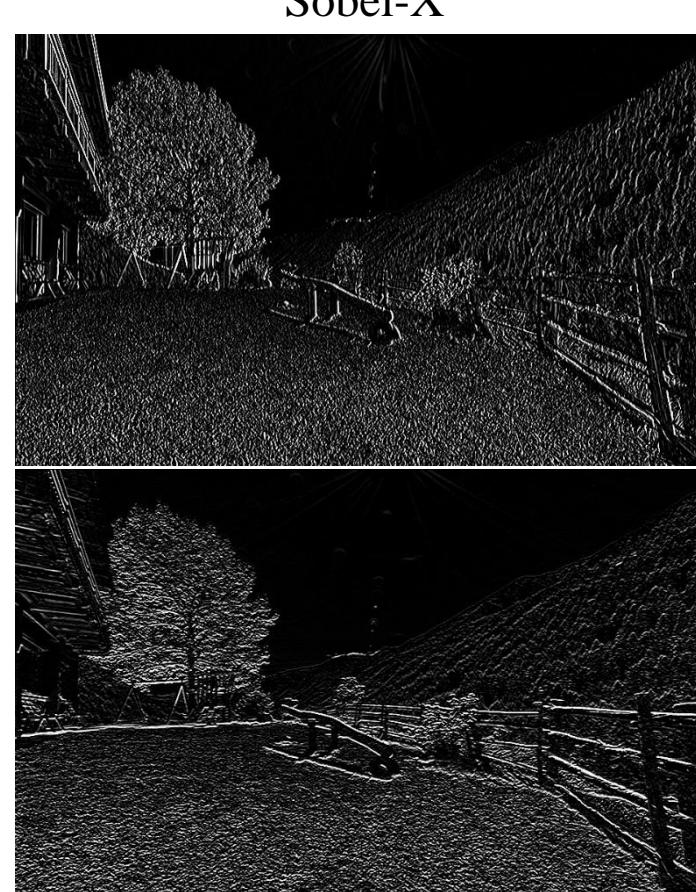


Derivative and Applications

Edge detection



Edge
detection



Sobel-Y

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * I$$

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} * I$$

Engineering Feature Extractors

❖ Edge detection

Tính đạo hàm trung bình theo hướng x

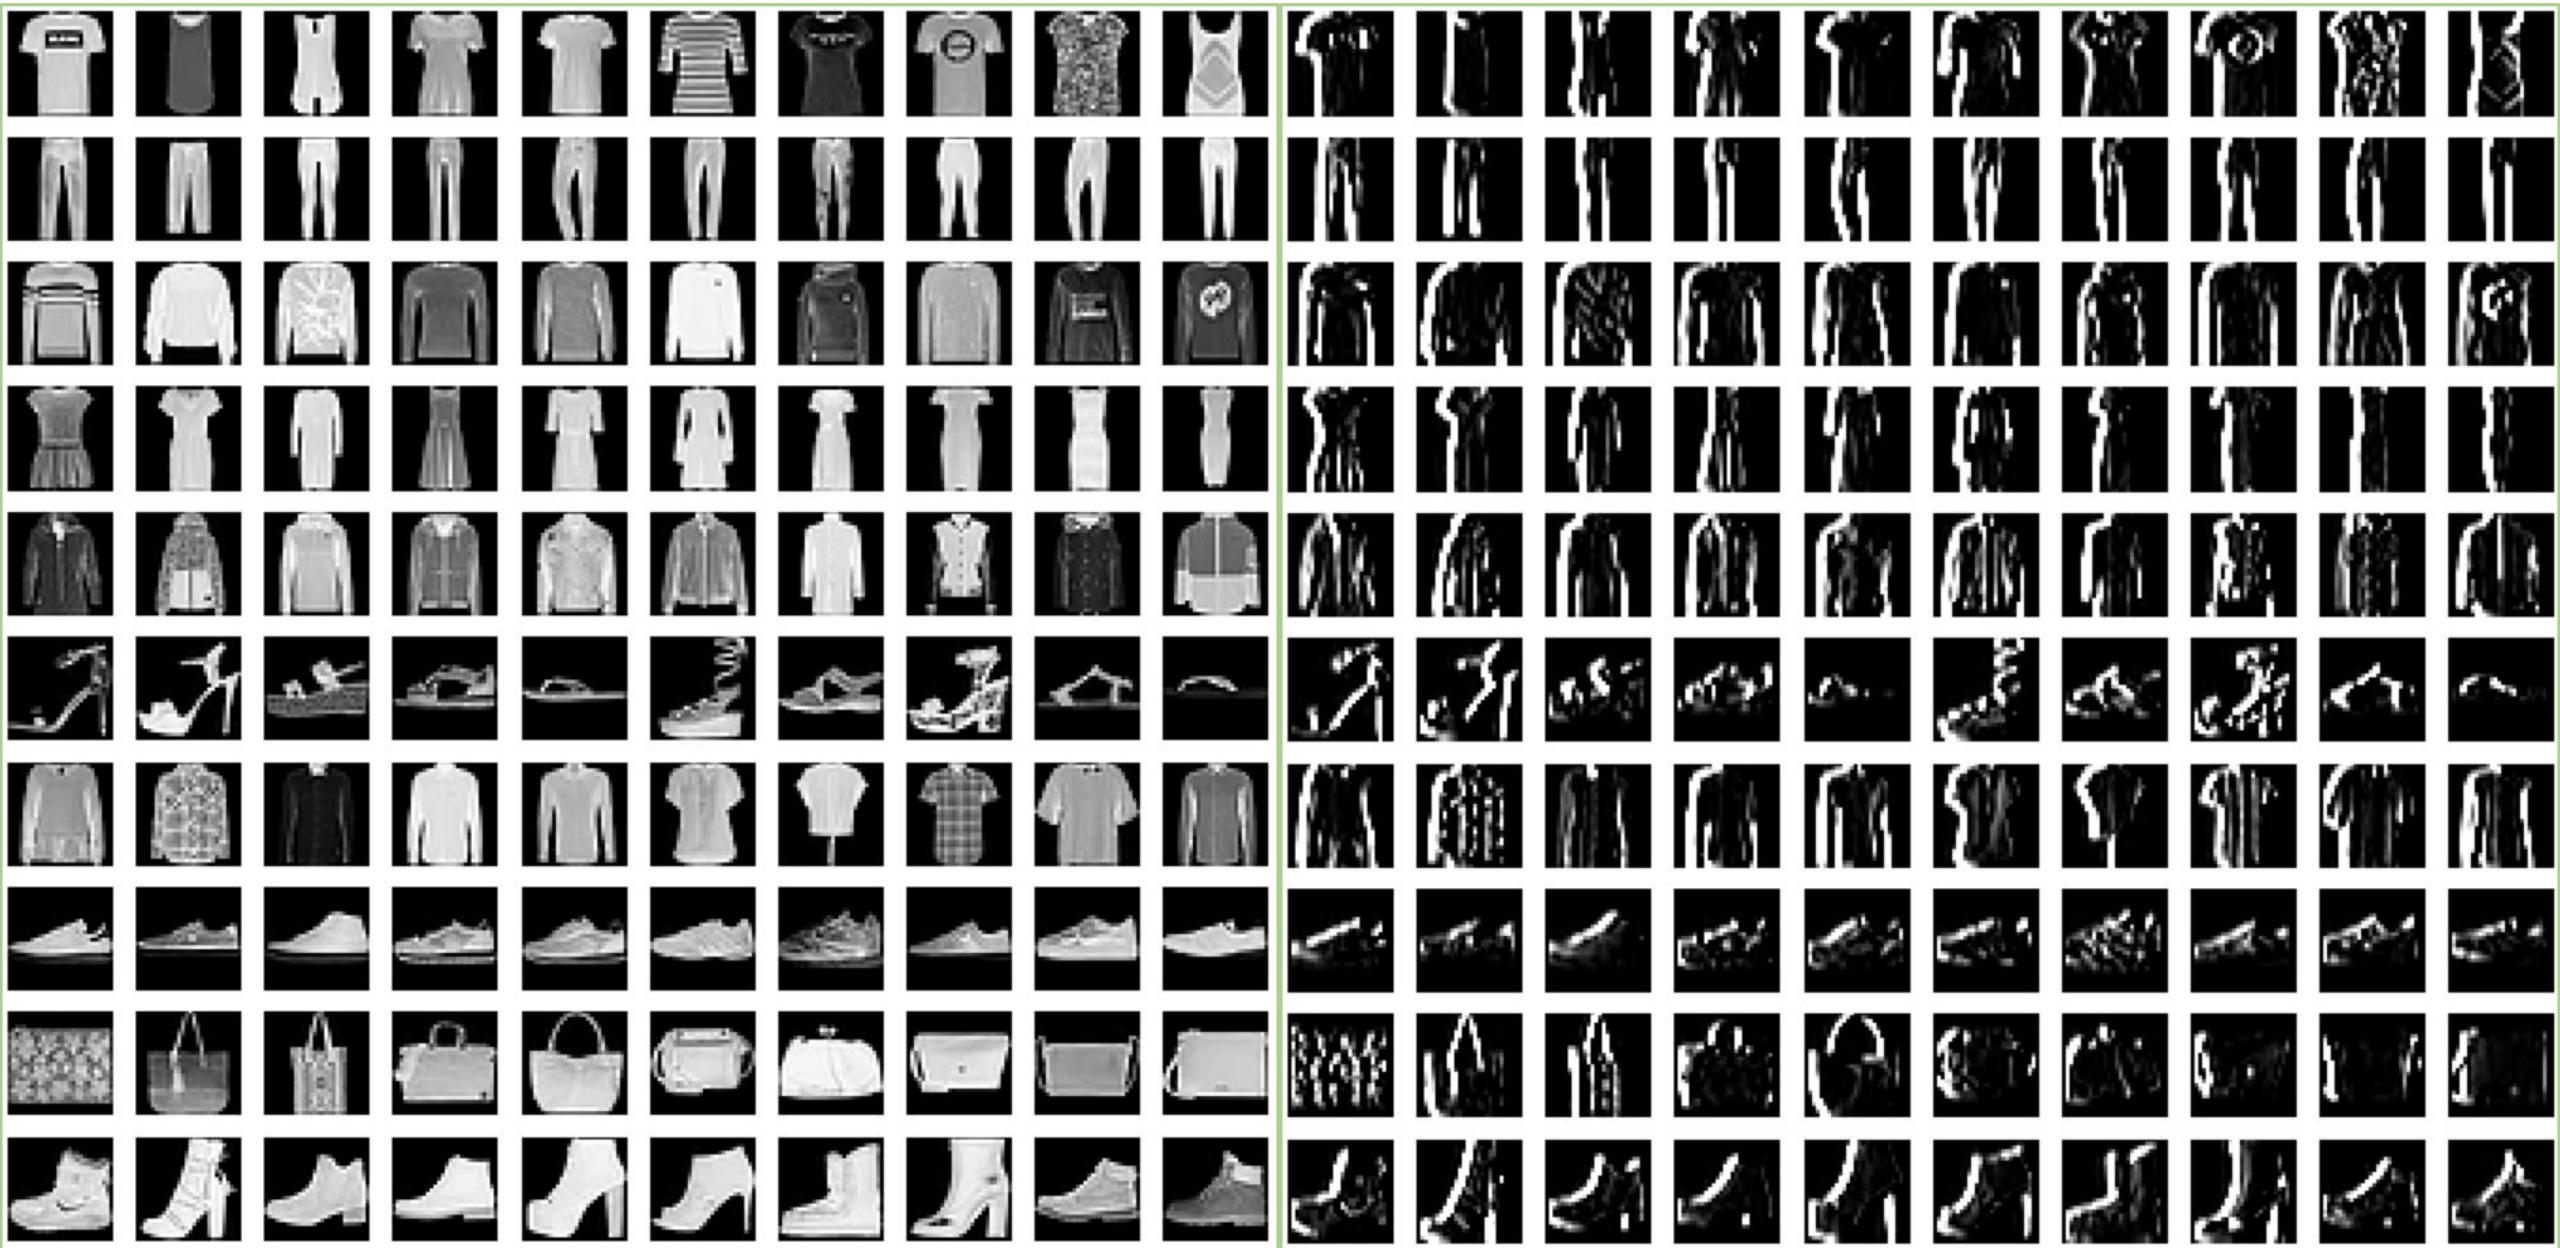
$$\begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \text{ weighted average} * \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array} \text{ x-derivative} = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array} \text{ Sobel for x direction}$$

```
@staticmethod
def _sobel(image, ksize=3):
    sobel_x = cv2.Sobel(image, cv2.CV_32F, 1, 0, ksize=ksize)
    sobel_y = cv2.Sobel(image, cv2.CV_32F, 0, 1, ksize=ksize)
    sobel_x = torch.from_numpy(sobel_x)
    sobel_y = torch.from_numpy(sobel_y)
    sobel_magnitude = torch.hypot(sobel_x, sobel_y)
    return sobel_magnitude
```

Tính đạo hàm trung bình theo hướng y

$$\begin{array}{|c|} \hline 1 \\ \hline 0 \\ \hline -1 \\ \hline \end{array} \text{ y-derivative} * \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array} \text{ weighted average} = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array} \text{ Sobel for y direction}$$

Sobel_X

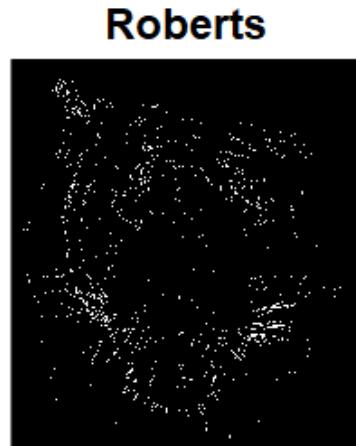
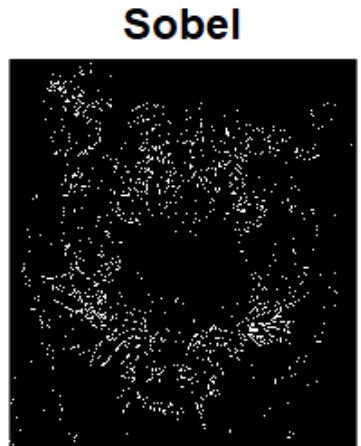
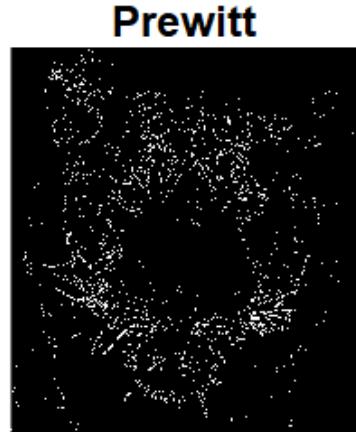
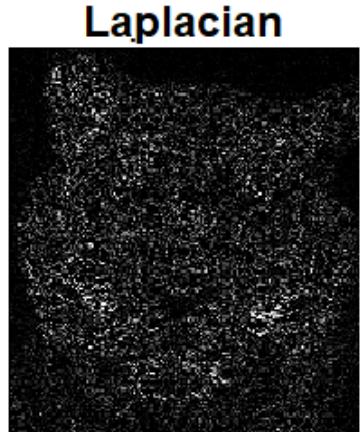


Sobel_Y



Engineering Feature Extractors

❖ Traditional methods



```
@staticmethod
def _scharr(image):
    scharr_x = cv2.Scharr(image, cv2.CV_32F, 1, 0)
    scharr_y = cv2.Scharr(image, cv2.CV_32F, 0, 1)
    scharr_x = torch.from_numpy(scharr_x)
    scharr_y = torch.from_numpy(scharr_y)
    scharr_magnitude = torch.hypot(scharr_x, scharr_y)
    return scharr_magnitude

@staticmethod
def _laplacian(image):
    laplacian_img = cv2.Laplacian(image, cv2.CV_32F)
    laplacian_img = torch.from_numpy(laplacian_img)
    return laplacian_img
```

Gradient_X

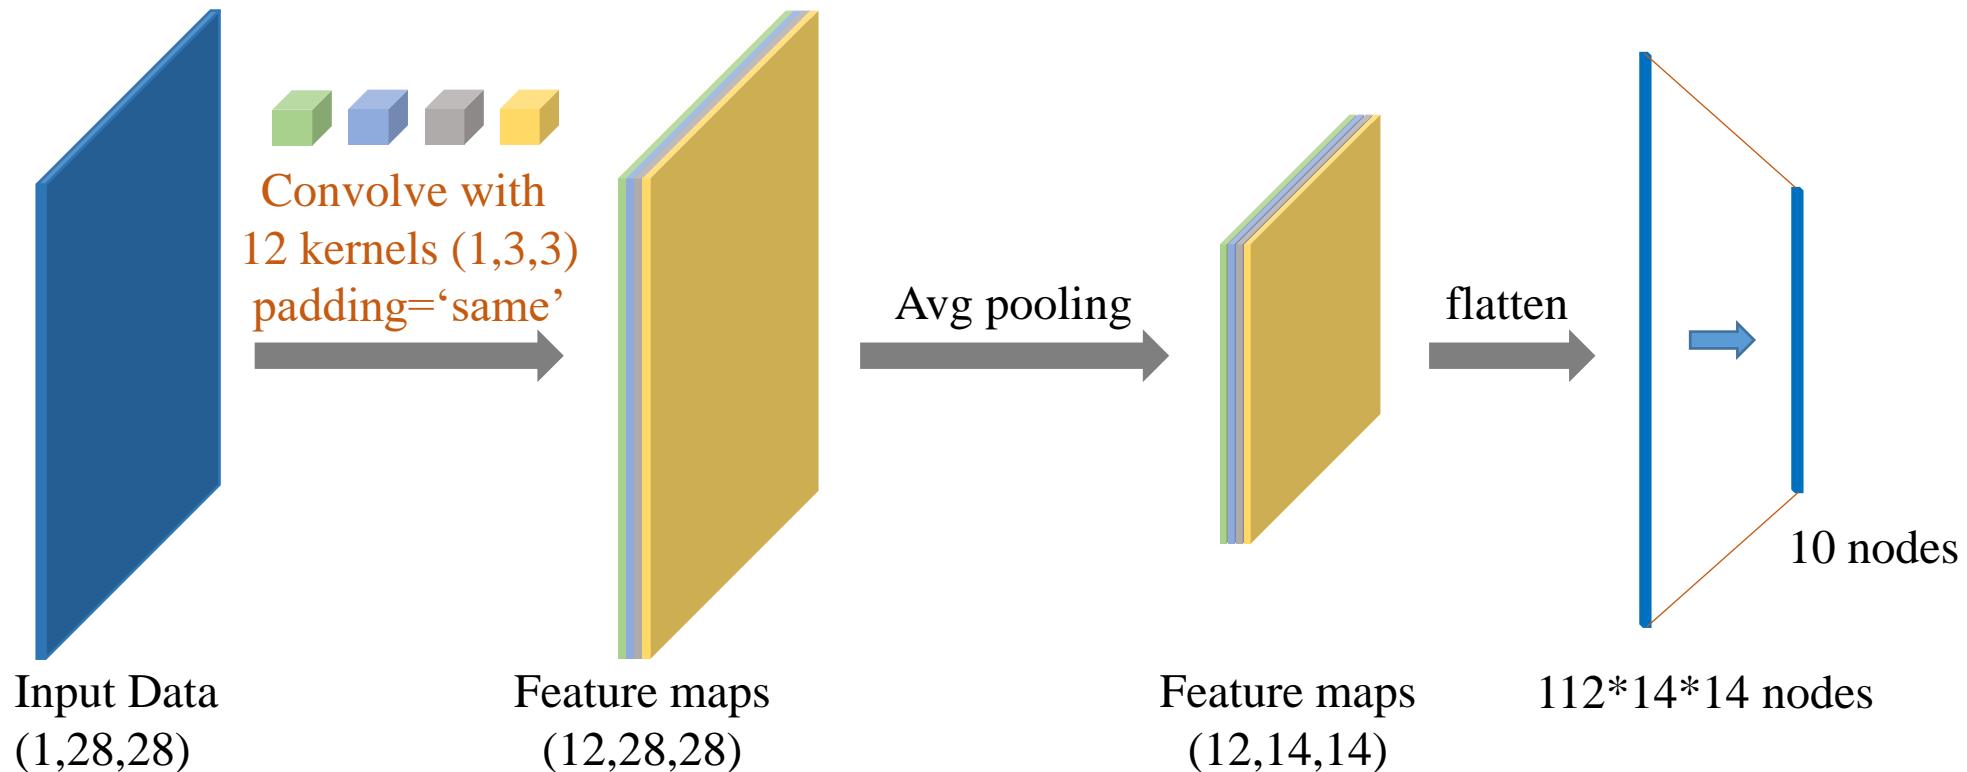


Gradient_Y



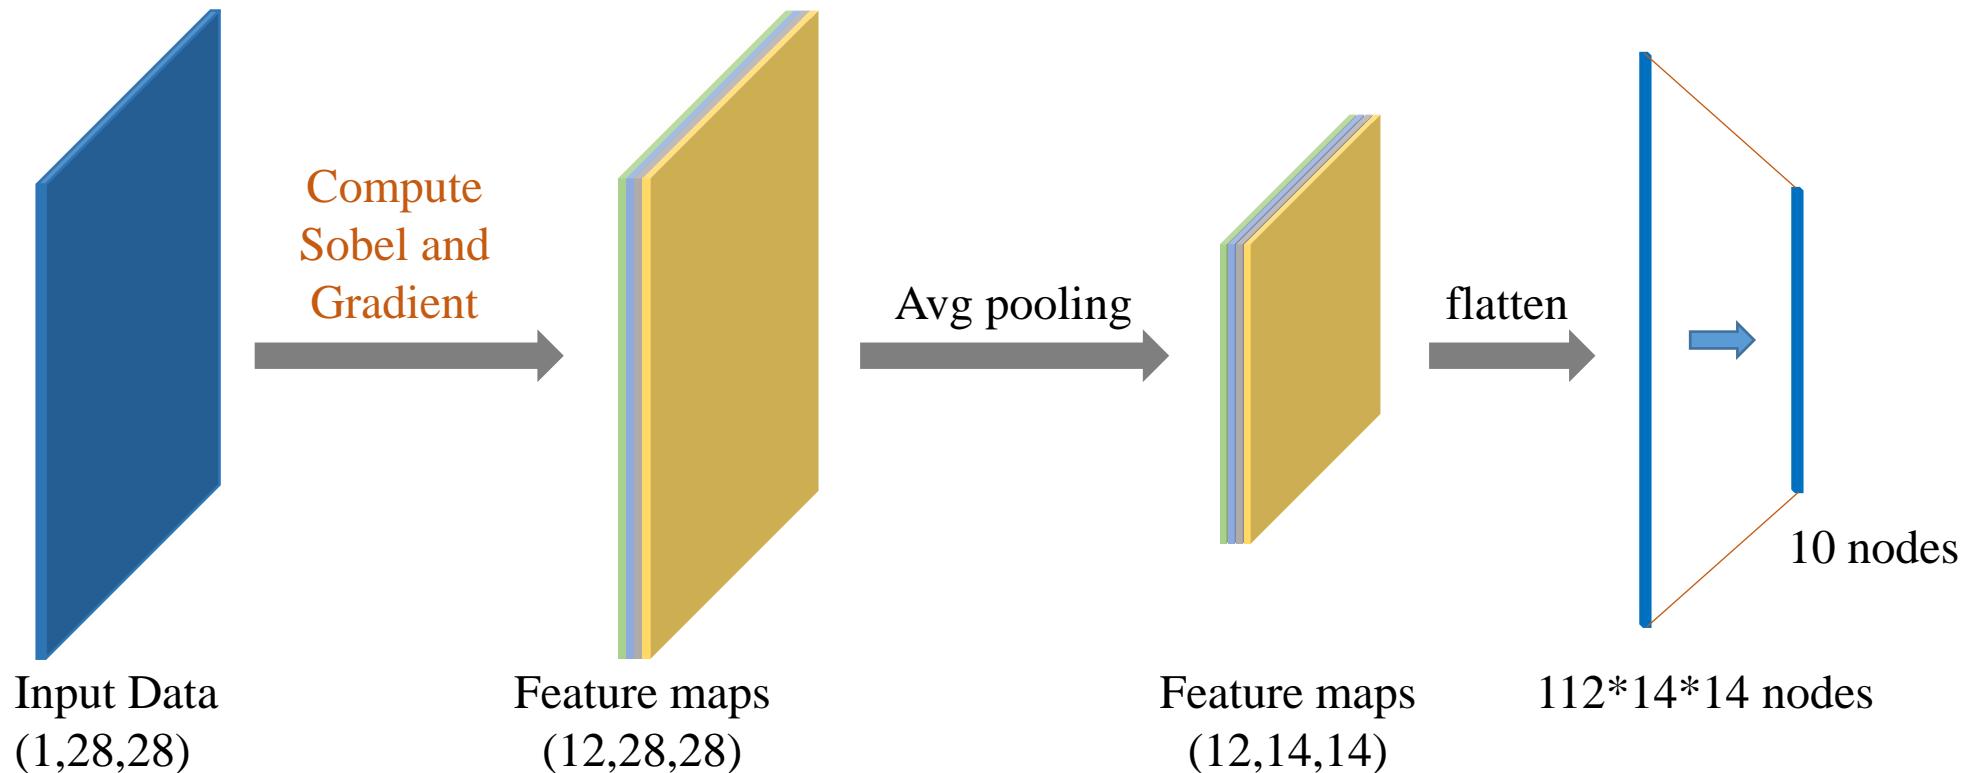
Comparison

❖ Engineering features vs. CNN features

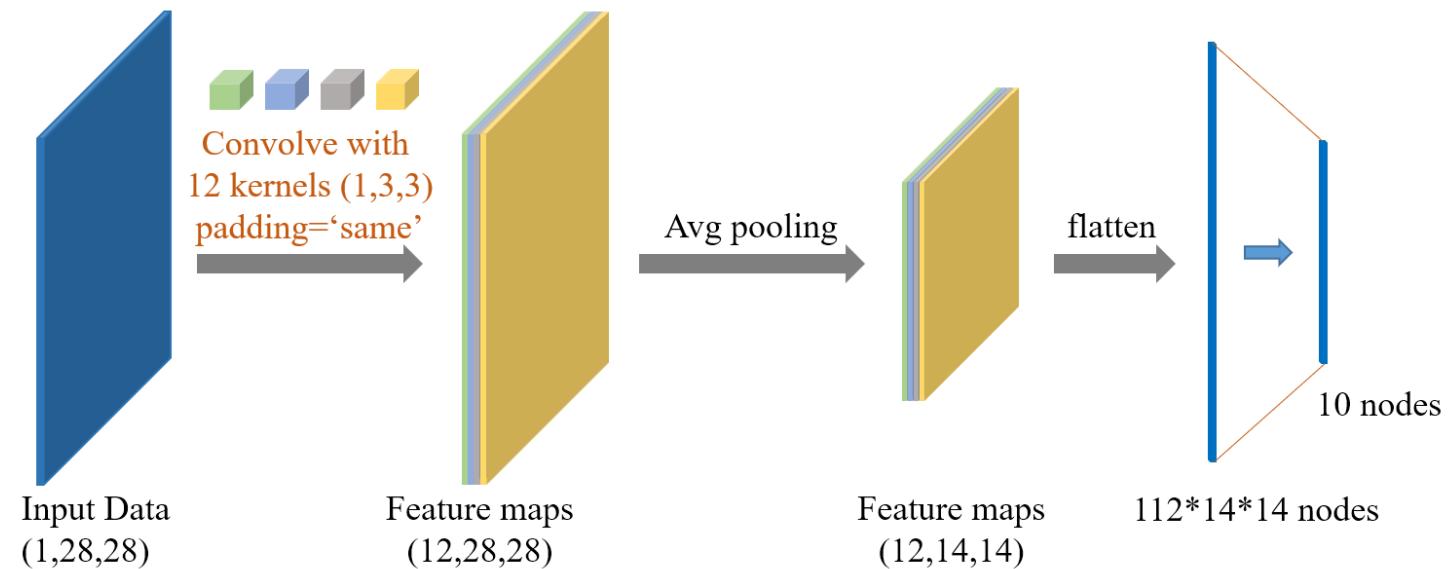
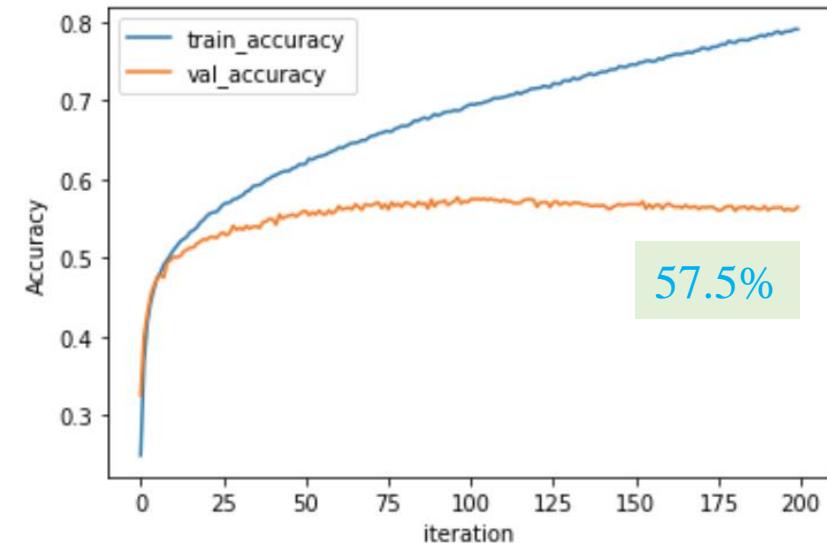
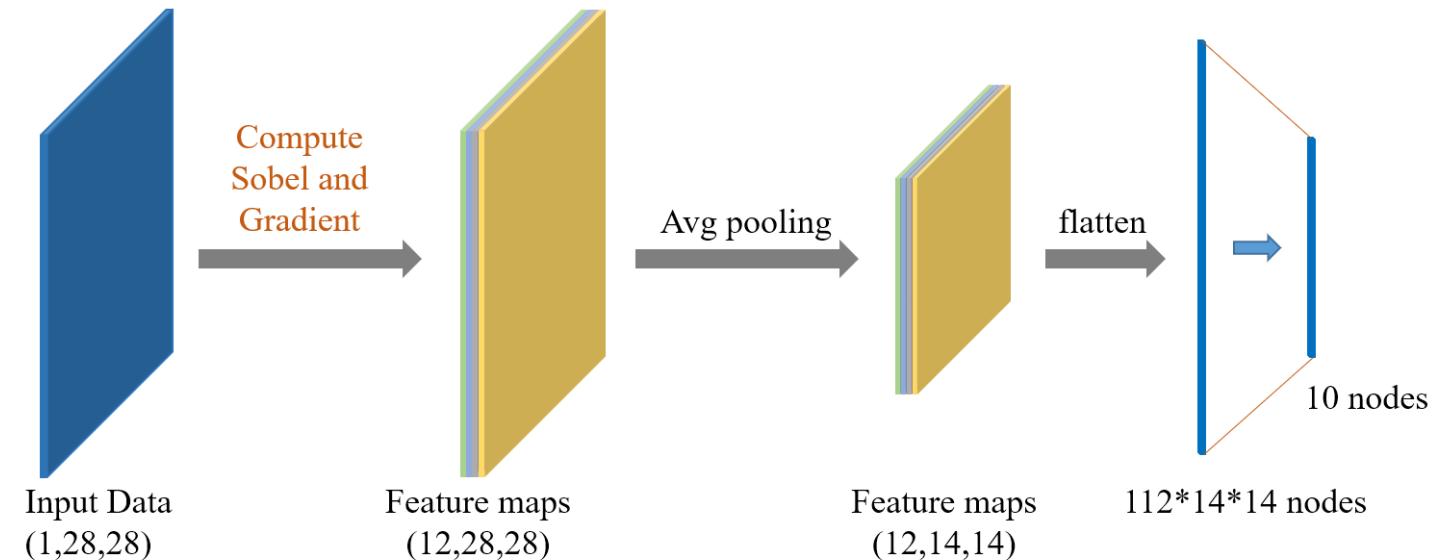
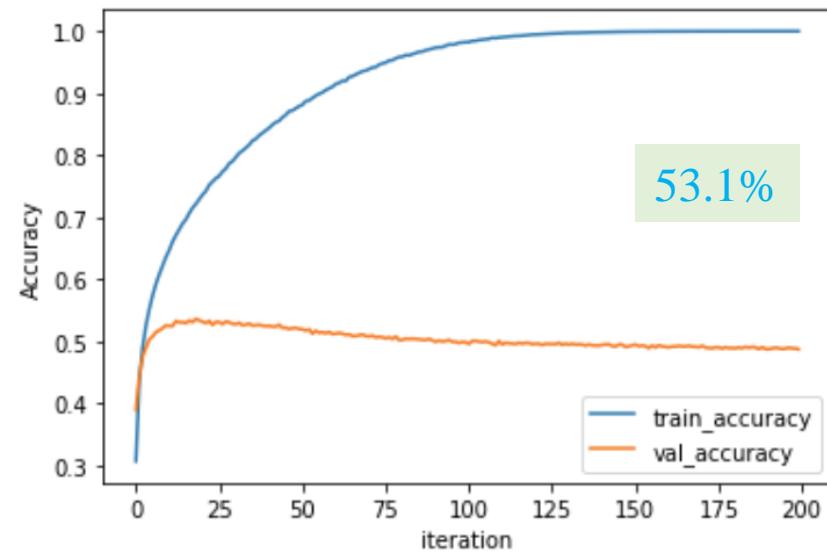


Comparison

❖ Engineering features vs. CNN features



Engineering Features vs. CNN Features using Cifar10



Outline

- Fashion-MNIST/Cifar10 Using only Conv2d
- Pooling
- Padding
- 1x1 Convolution
- LeNet and VGG Models

1x1 Convolution

- ❖ Why 1x1 Convolution
 - ❖ Flexible input size



Yann LeCun

April 7, 2015 ·

...

In Convolutional Nets, there is no such thing as "fully-connected layers".
There are only convolution layers with 1x1 convolution kernels and a full connection table.

It's a too-rarely-understood fact that ConvNets don't need to have a fixed-size input. You can train them on inputs that happen to produce a single output vector (with no spatial extent), and then apply them to larger images. Instead of a single output vector, you then get a spatial map of output vectors. Each vector sees input windows at different locations on the input.

In that scenario, the "fully connected layers" really act as 1x1 convolutions.

Yann LeCun

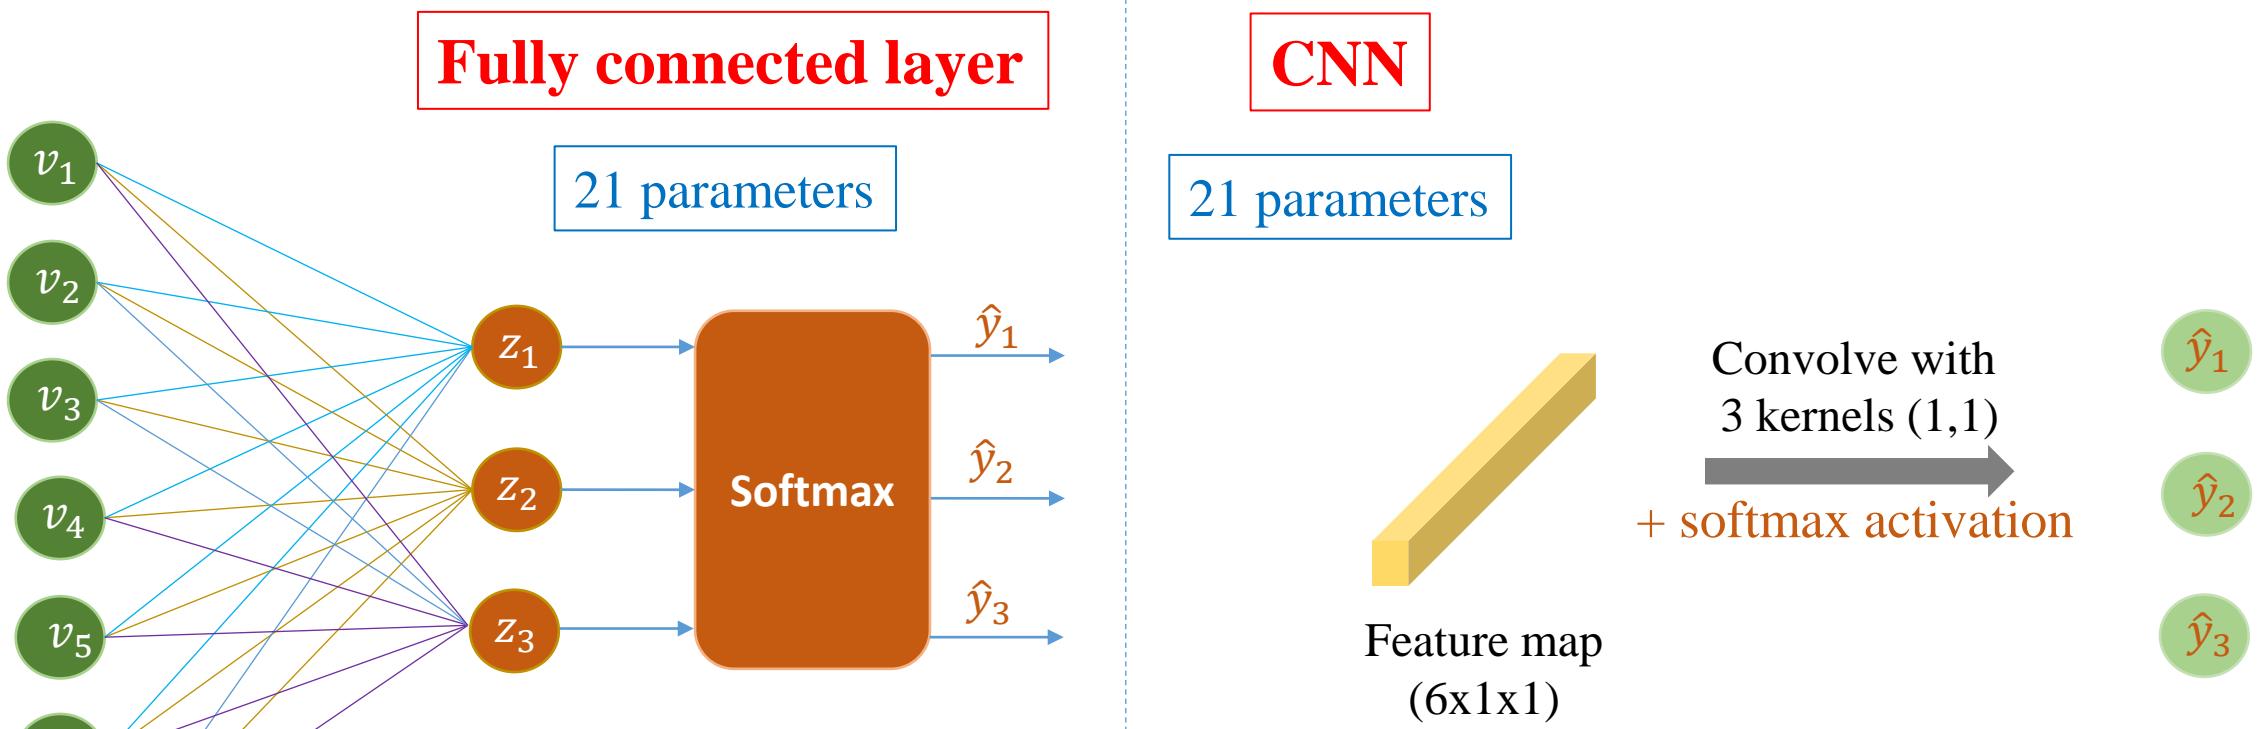


Yann LeCun in 2018

Born	July 8, 1960 (age 60) Soisy-sous-Montmorency, France
Alma mater	ESIEE Paris (MSc) Pierre and Marie Curie University (PhD)
Known for	Deep learning
Awards	Turing Award (2018) AAAI Fellow (2019) Legion of Honour (2020)
	Scientific career
Institutions	Bell Labs (1988-1996) New York University Facebook
Thesis	<i>Modèles connexionnistes de l'apprentissage (connectionist learning models)</i> (1987a)
Doctoral advisor	Maurice Milgram
Website	yann.lecun.com

1x1 Convolution

❖ Comparison



1x1 Convolution

Replace FC by 1x1 Conv

(3x3) Conv with stride=1,
padding='same' + ReLU



(3x3) Conv with stride=2,
padding='same' + ReLU



(7x7) Conv + ReLU

(1,28,28)



(32,28,28)



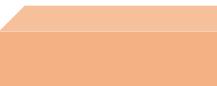
(32,14,14)



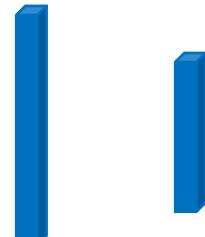
(64,14,14)



(64,7,7)



(64*7*7,)



(10,)



```
LeNet = nn.Sequential(  
    nn.Conv2d(1, 32, kernel_size=3, padding=1), nn.ReLU(), nn.MaxPool2d(2, 2),  
    nn.Conv2d(32, 64, kernel_size=3, padding=1), nn.ReLU(), nn.MaxPool2d(2, 2),  
    nn.Flatten(),  
    nn.Linear(7*7*64, 128),  
    nn.ReLU(),  
    nn.Linear(128, 10)  
)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 28, 28]	320
ReLU-2	[-1, 32, 28, 28]	0
MaxPool2d-3	[-1, 32, 14, 14]	0
Conv2d-4	[-1, 64, 14, 14]	18,496
ReLU-5	[-1, 64, 14, 14]	0
MaxPool2d-6	[-1, 64, 7, 7]	0
Flatten-7	[-1, 3136]	0
Linear-8	[-1, 128]	401,536
ReLU-9	[-1, 128]	0
Linear-10	[-1, 10]	1,290

1x1 Convolution

Replace FC by 1x1 Conv



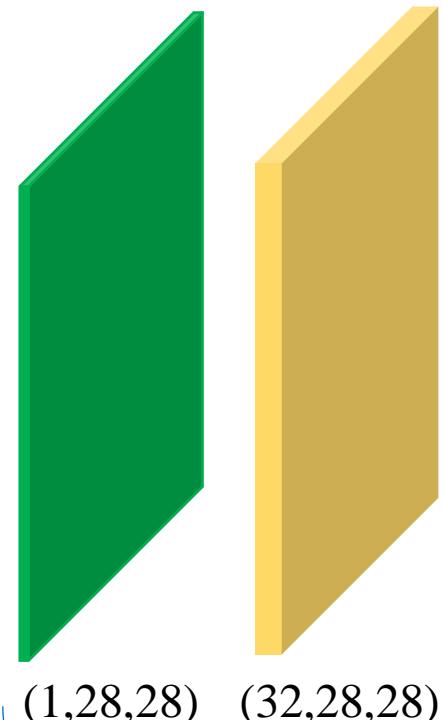
(3x3) Conv with stride=1,
padding='same' + ReLU



(3x3) Conv with stride=2,
padding='same' + ReLU



(7x7) Conv + ReLU



(32,14,14)

(64,14,14)

Feature extraction

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 28, 28]	320
ReLU-2	[-1, 32, 28, 28]	0
MaxPool2d-3	[-1, 32, 14, 14]	0
Conv2d-4	[-1, 64, 14, 14]	18,496
ReLU-5	[-1, 64, 14, 14]	0
MaxPool2d-6	[-1, 64, 7, 7]	0
Conv2d-7	[-1, 128, 1, 1]	401,536
ReLU-8	[-1, 128, 1, 1]	0
Conv2d-9	[-1, 10, 1, 1]	1,290
Flatten-10	[-1, 10]	0

(64,7,7)

(128,1,1)

(10,1,1)
(10,)

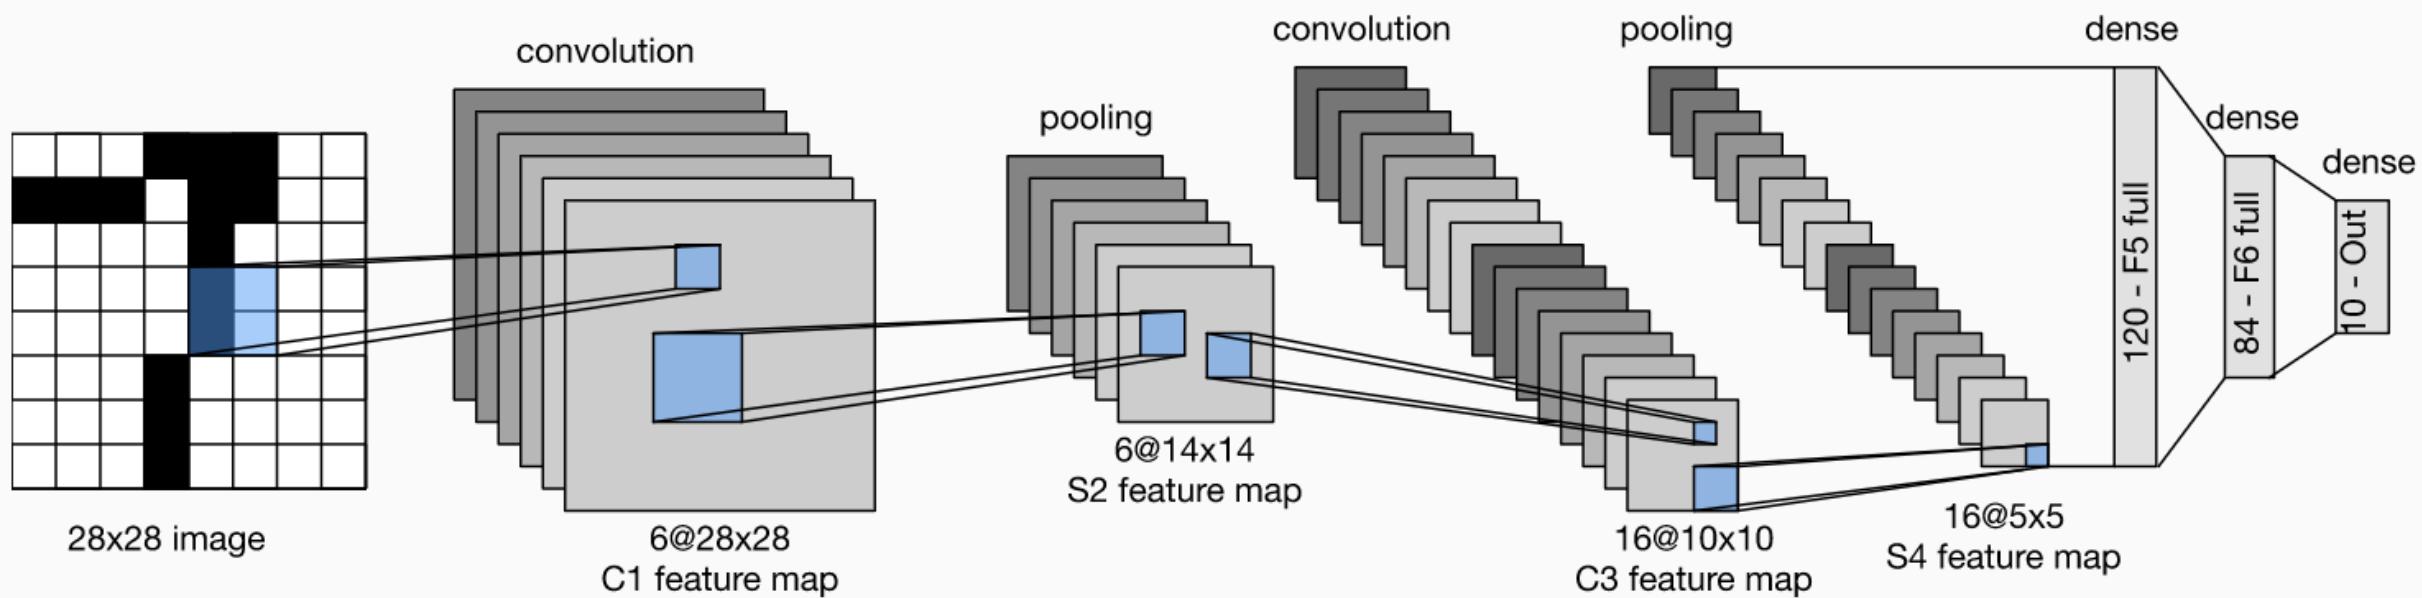
Classification

Outline

- Fashion-MNIST/Cifar10 Using only Conv2d
- Pooling
- Padding
- 1x1 Convolution
- LeNet and VGG Models

LeNet Architecture

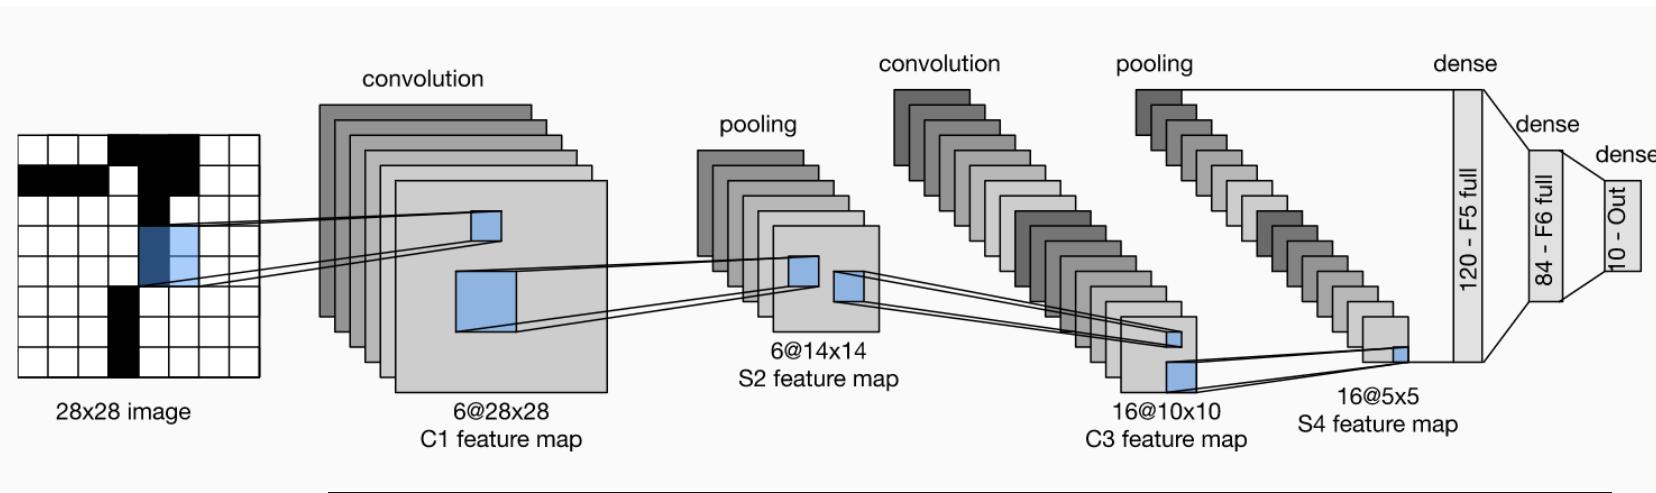
❖ Model Construction



LeNet Architecture

❖ Model Construction

```
LeNet = nn.Sequential(  
    nn.Conv2d(1, 6, 5, padding=2),  
    nn.ReLU(),  
    nn.MaxPool2d(2, 2),  
  
    nn.Conv2d(6, 16, 5),  
    nn.ReLU(),  
    nn.MaxPool2d(2, 2),  
  
    nn.Flatten(),  
    nn.Linear(16*5*5, 120),  
    nn.ReLU(),  
  
    nn.Linear(120, 84),  
    nn.ReLU(),  
    nn.Linear(84, 10)  
)
```

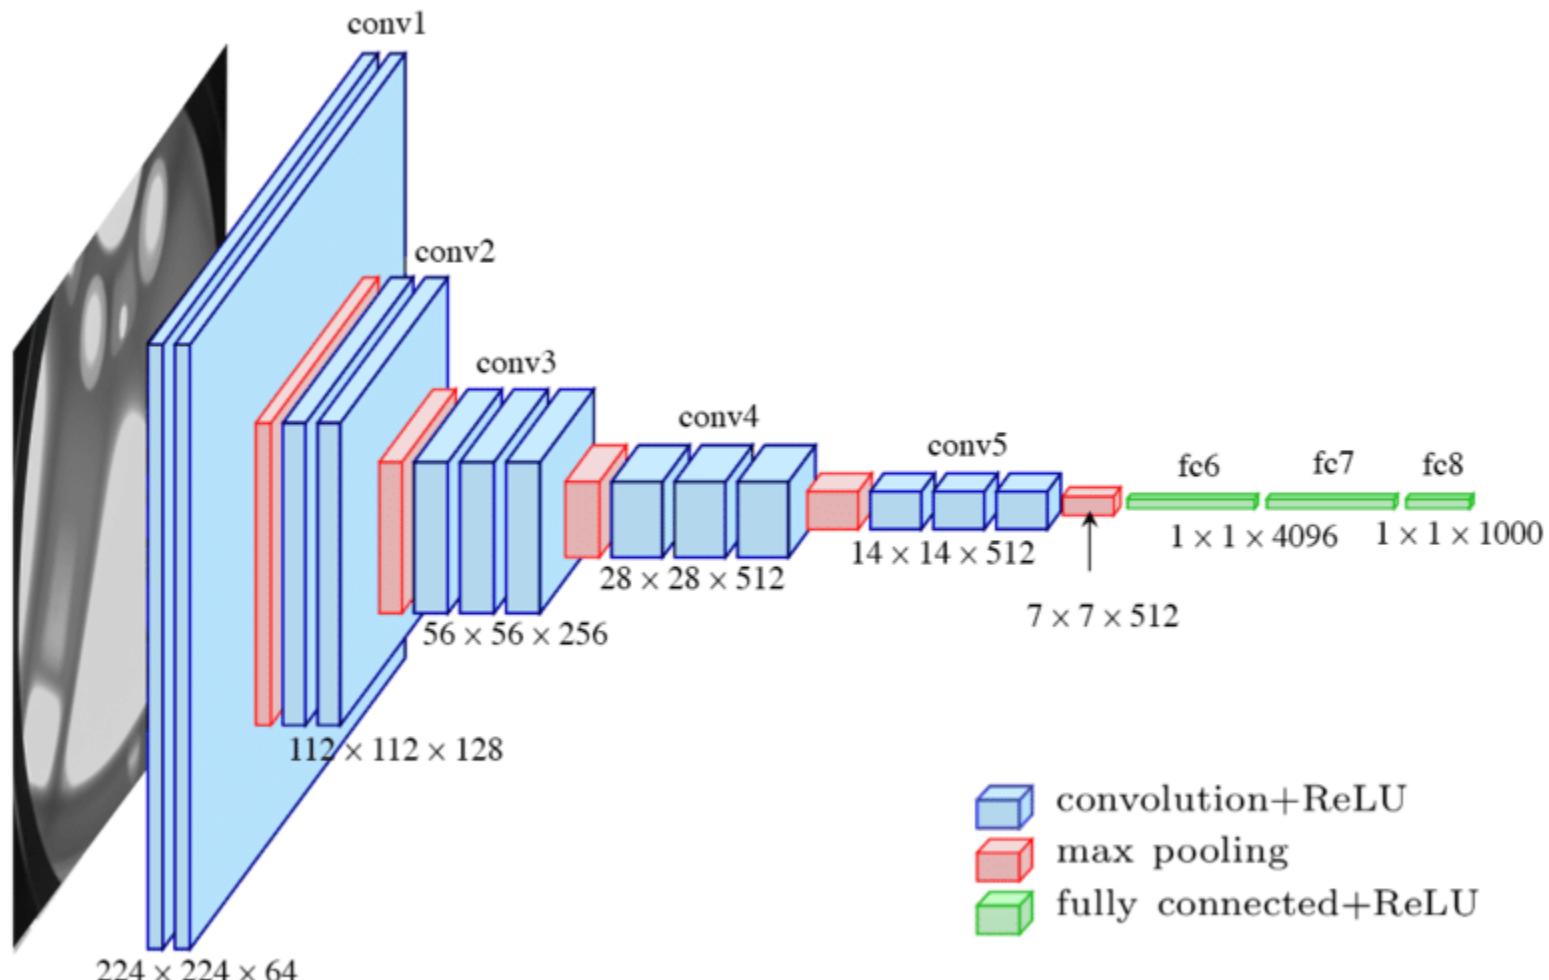


Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 6, 28, 28]	156
ReLU-2	[-1, 6, 28, 28]	0
MaxPool2d-3	[-1, 6, 14, 14]	0
Conv2d-4	[-1, 16, 10, 10]	2,416
ReLU-5	[-1, 16, 10, 10]	0
MaxPool2d-6	[-1, 16, 5, 5]	0
Flatten-7	[-1, 400]	0
Linear-8	[-1, 120]	48,120
ReLU-9	[-1, 120]	0
Linear-10	[-1, 84]	10,164
ReLU-11	[-1, 84]	0
Linear-12	[-1, 10]	850

Total params: 61,706
Trainable params: 61,706
Non-trainable params: 0

VGG16 Architecture

❖ Model Construction



```

self.relu = nn.ReLU()
self.features = nn.Sequential(
    nn.Conv2d(3, 64, kernel_size=3, padding=1), self.relu,
    nn.Conv2d(64, 64, kernel_size=3, padding=1), self.relu,
    nn.MaxPool2d(kernel_size=2, stride=2),

    nn.Conv2d(64, 128, kernel_size=3, padding=1), self.relu,
    nn.Conv2d(128, 128, kernel_size=3, padding=1), self.relu,
    nn.MaxPool2d(kernel_size=2, stride=2),

    nn.Conv2d(128, 256, kernel_size=3, padding=1), self.relu,
    nn.Conv2d(256, 256, kernel_size=3, padding=1), self.relu,
    nn.Conv2d(256, 256, kernel_size=3, padding=1), self.relu,
    nn.MaxPool2d(kernel_size=2, stride=2),

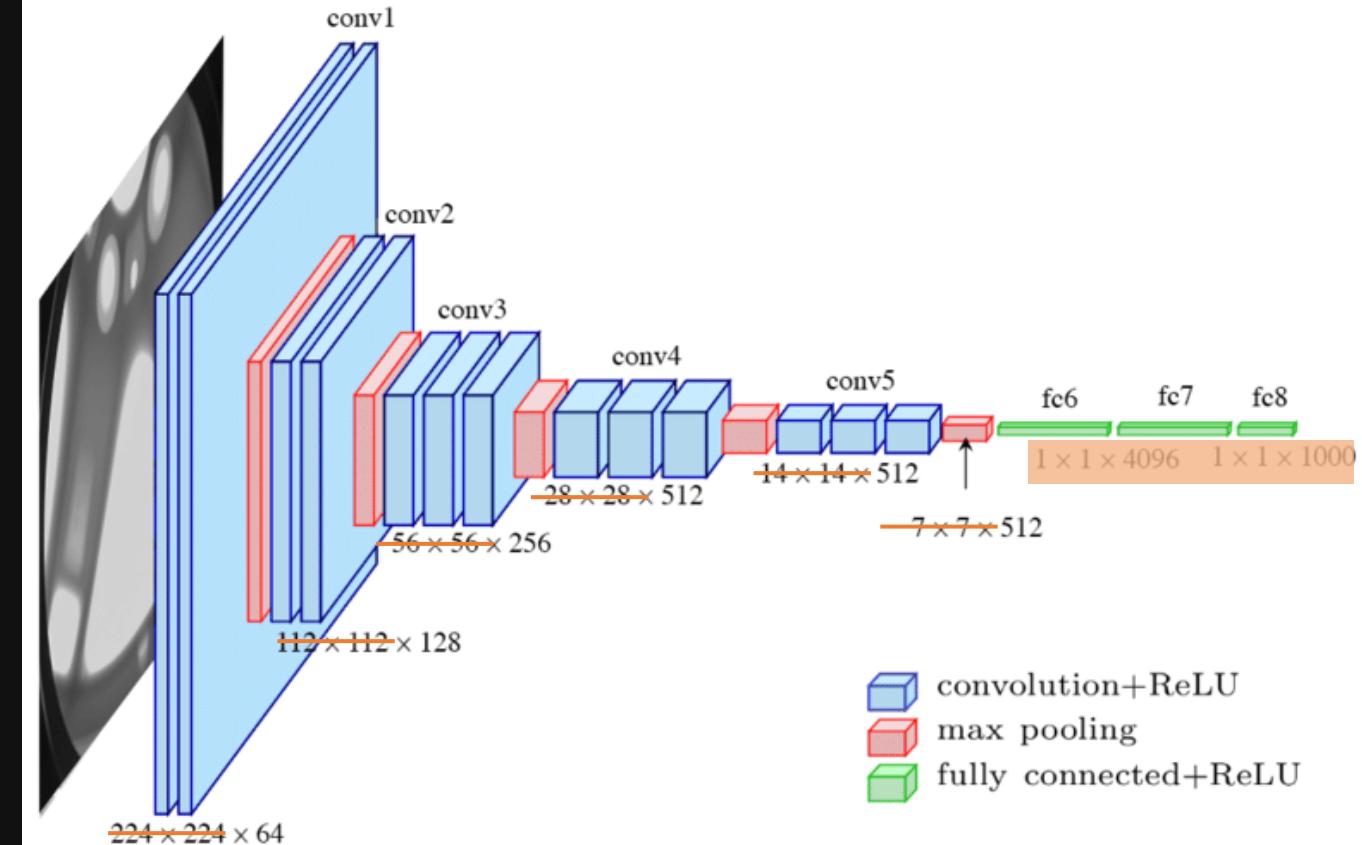
    nn.Conv2d(256, 512, kernel_size=3, padding=1), self.relu,
    nn.Conv2d(512, 512, kernel_size=3, padding=1), self.relu,
    nn.Conv2d(512, 512, kernel_size=3, padding=1), self.relu,
    nn.MaxPool2d(kernel_size=2, stride=2),
)

self.classifier = nn.Sequential(
    nn.Flatten(), nn.Linear(512 * 7 * 7, 128), self.relu,
    nn.Linear(128, 128), self.relu, nn.Linear(128, 10)
)

```

VGG16 Architecture

❖ Model Construction



Further Reading

❖ Reading

<https://cs231n.github.io/convolutional-networks/>

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

<https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-convolutional-neural-networks>

